

# Events

---

- The cornerstone of event driven architecture
- Require a well-defined definition
- Evolved from other architectures

# Microservices Architecture


---

- Based on loosely-coupled services
- Each service in its own process
- Lightweight communication protocols
- Polyglot
  - No platform dependency between services
- Replaces two legacy architectures



# Monolith






**SOA**


**(Service Oriented Architecture)**

# Typical Microservices System



# Microservices Communication

- Perhaps the most important part in microservices architecture
- Dictates performance, scalability, implementation and more
- Event Driven Architecture handles the communication part



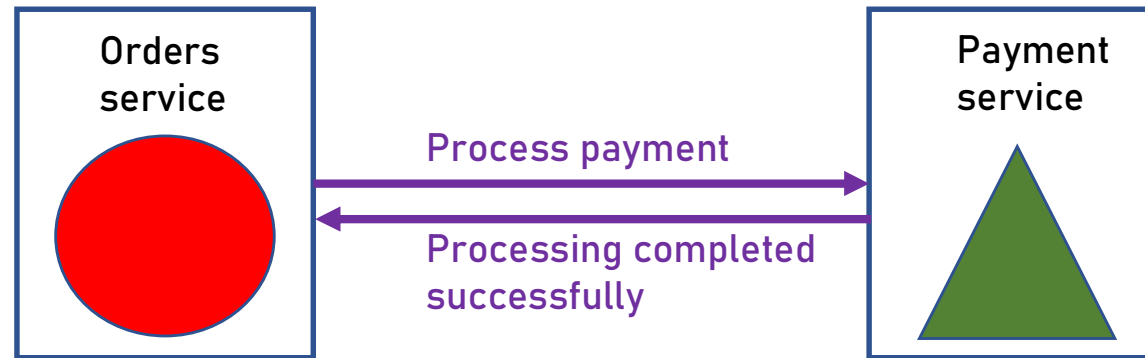
# Command and Query

---

- The classic communication between services
- Services either:
  - Send command
  - Query for data

# Command

- Service asks another service to do something




- There might be a response to the command, usually a success or failure indicator



# Query

- Service asks another service for data



- There's always a response to the query, containing the data

# Command and Query

- Main characteristics:

## Command

- Do something
- Usually synchronous
- Sometimes returns a response
- Calling service needs to know who handles the command

## Query

- Retrieve data
- Almost always synchronous
- Always returns a response
- Calling service needs to know who handles the query

# Problems with Command and Query

---

- Three major problems with command and query:

Performance

Coupling

Scalability

# Performance

## Command

- Do something
- Usually synchronous ←
- Sometimes returns a response
- **Calling service needs to know who handles the command**

## Query

- Retrieve data
- Almost always synchronous ←
- Always returns a response
- **Calling service needs to know who handles the query**

- Synchronous = the calling service waits for the command / query to complete
- Potential for performance hit

# Coupling

## Command

- Do something
- Usually synchronous
- Sometimes returns a response
- **Calling service needs to know who handles the command** ←

## Query

- Retrieve data
- Almost always synchronous
- Always returns a response
- **Calling service needs to know who handles the query** ←

- The calling service calls a specific service
- If the called service changes – the calling service has to change too
- More work, more maintenance



# Scalability

## Command

- Do something
- Usually synchronous
- Sometimes returns a response
- **Calling service needs to know who handles the command** ←

## Query

- Retrieve data
- Almost always synchronous
- Always returns a response
- **Calling service needs to know who handles the query** ←

- The calling service calls a single instance of a service
- If this instance is busy – there's a performance hit
- Adding another instances is possible, but difficult
  - Add load balancer, configure probes etc.

# Event

- Indicates that something happened in the system



- There's never a response to the event

# Event

---

- Main characteristics:

## Event

- Something happened
- Asynchronous
- Never returns a response
- Calling service has no idea who handles the event

# Contents of Event

- Two types of event data:

## Complete

- Contains all the relevant data
- Usually entity data
- No additional data is required for the event processing
- Example:  
event\_type: CustomerCreated  
customer\_id: 17  
first\_name: David  
last\_name: Jones  
join\_date: 2022-03-15

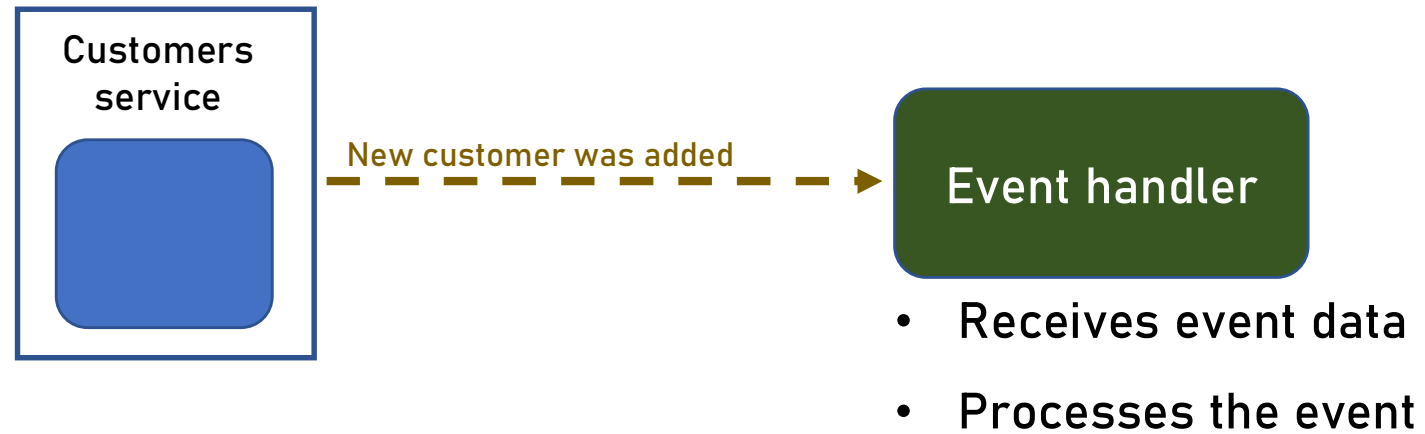
## Pointer

- Contains pointer to the complete data of the entity
- Complete data usually stored in a database
- Event handler needs to access the database to retrieve complete data
- Example:  
event\_type: CustomerCreated  
customer\_id: 17



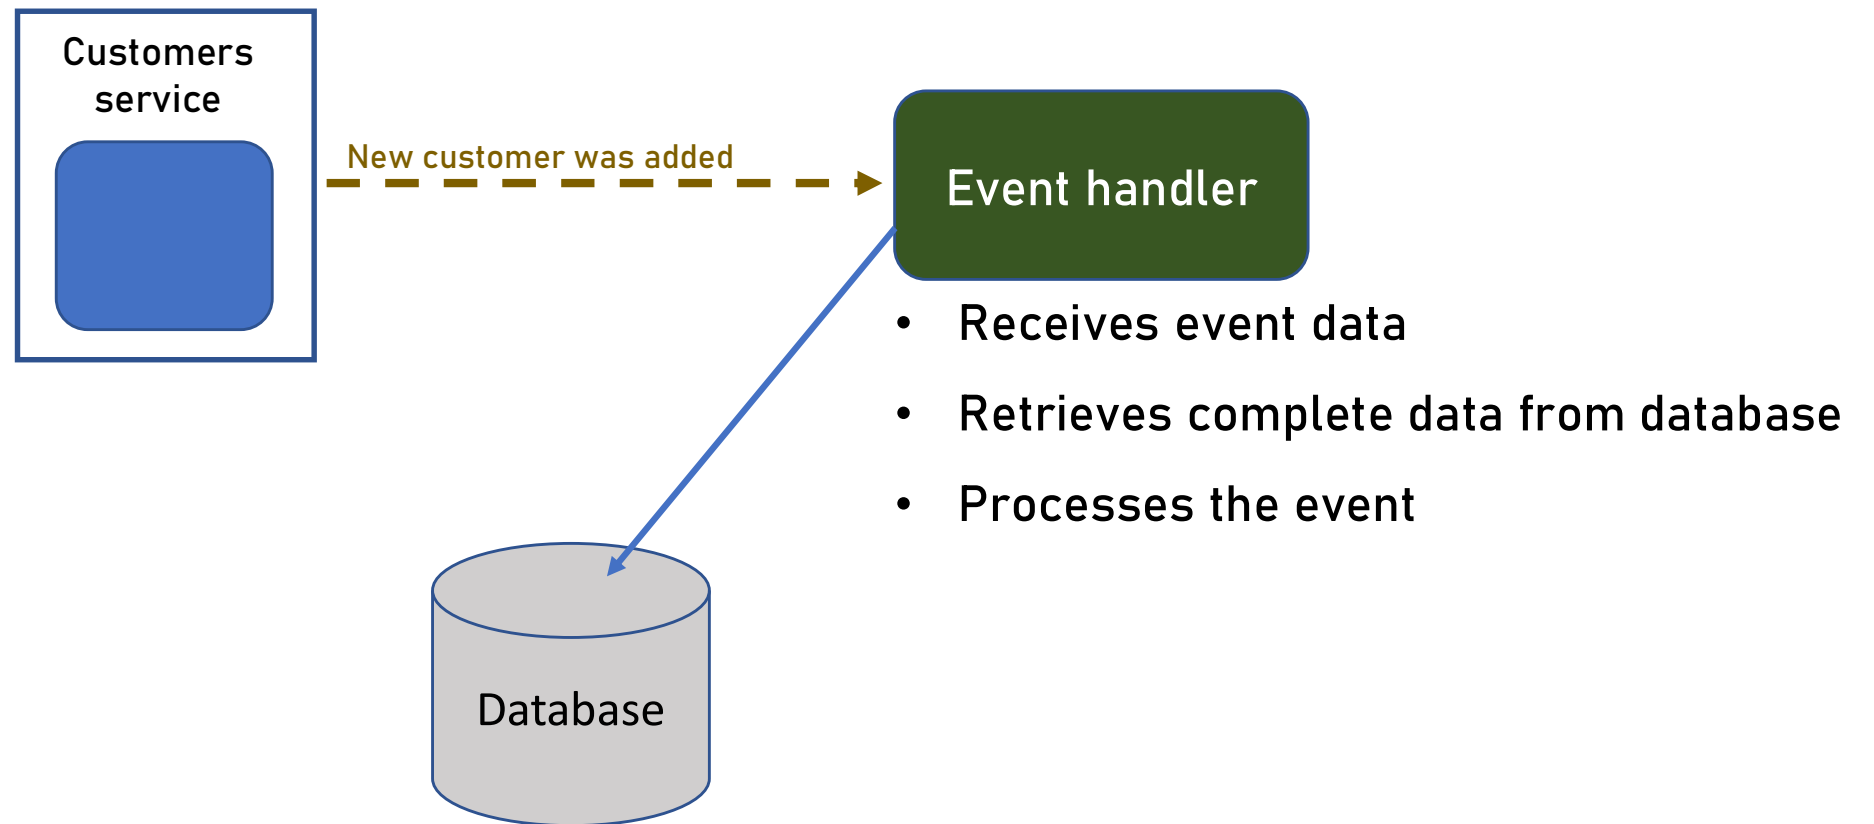
Pointer

# Flow of Complete Event Handling





# Flow of Pointer Event Handling



# Complete vs Pointer

- When to use which?

## Complete

- The better approach
- Makes the event completely autonomous
- Can get out of the system boundaries

## Pointer


- Use when:
  - Data is large
  - Need to ensure data is up-to-date
    - Assuming database is a single-source-of-truth

# Event Driven Architecture


---

- A software architecture paradigm that uses events as the mean of communication between services
- Often called EDA
- Has three main components

# Event Driven Architecture Components



# Producer



Producer



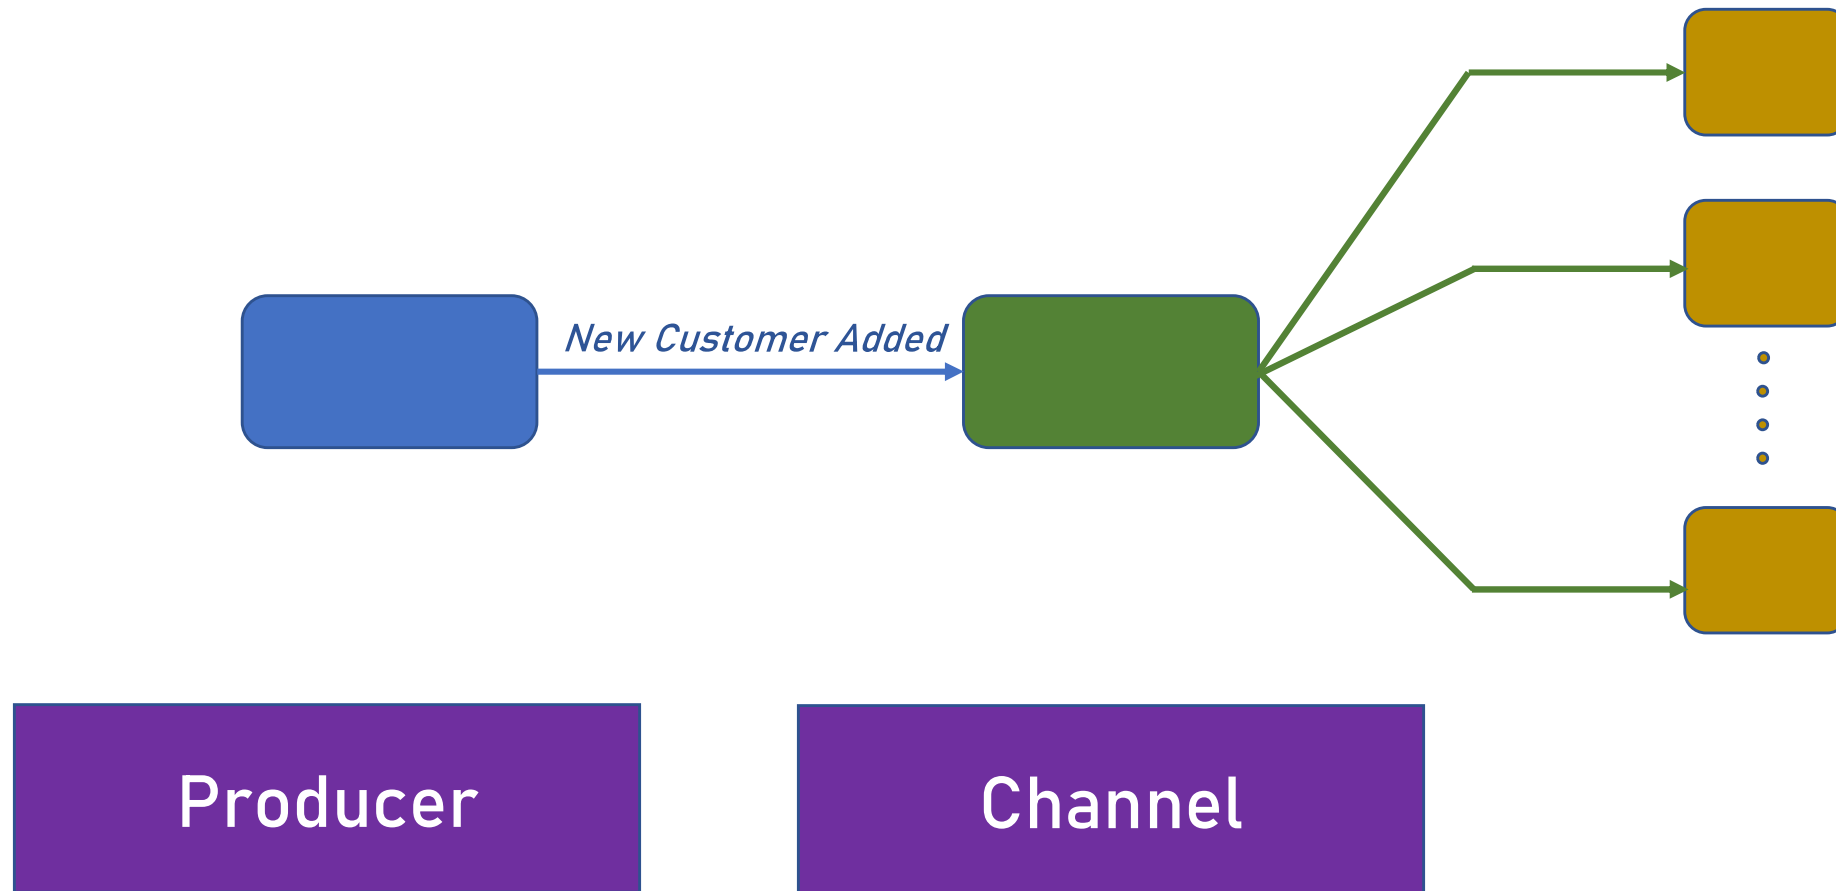
# Producer

---

- The component / service sending the event
- Often called Publisher
- Usually sends event reporting something the component done
- Examples:
  - Customer service -> *New Customer Added* event
  - Inventory service -> *Item Sold Out* event

# Producer

- The producer sends the event to the Channel




# Producer

---

- Exact method of calling the channel depends on the channel
- Usually using a dedicated SDK developed by the channel vendor
- Utilizes some kind of network call, usually with specialized ports and proprietary protocol
- I.e.: RabbitMQ listens on port 5672 and uses the AMQP protocol
- Producer can be developed using any development language

# Channel



Channel

# Channel

---

- The most important component in the Event Driven Architecture
- Responsible for distributing the events to the relevant parties
- The channel places the event in a specialized queue, often called  
Topic or Fanout
- Consumers listen to this queue and grab the event



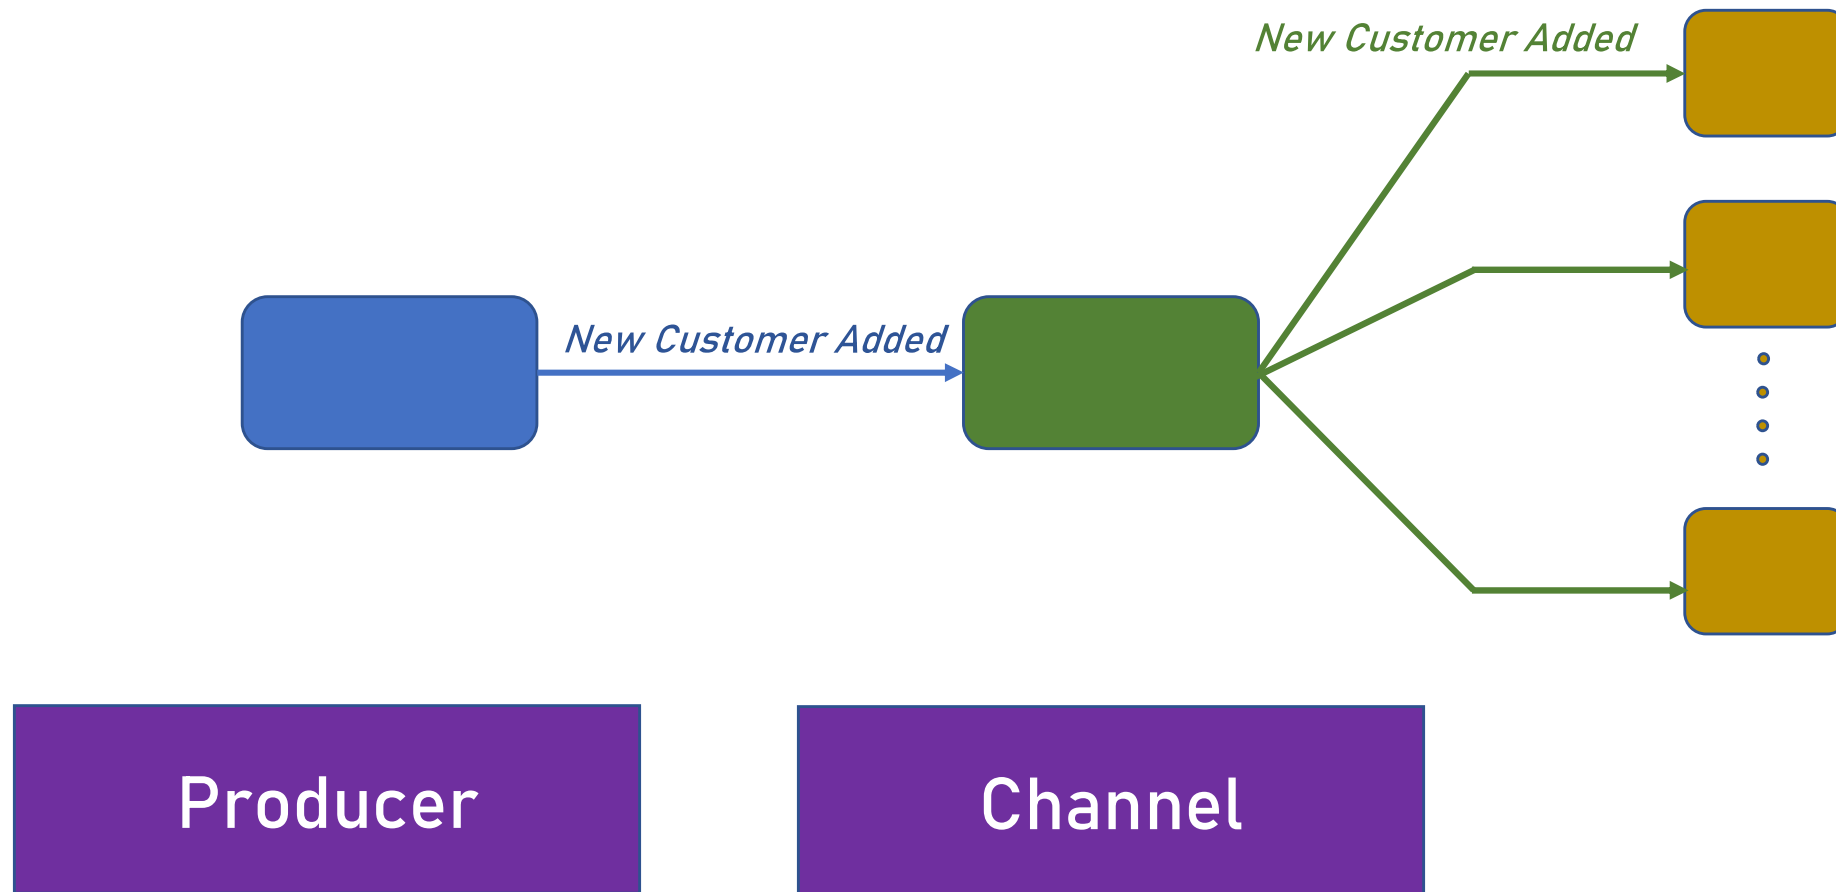
# Channel

---

- Note:
  - Implementation details vary wildly between channels
  - RabbitMQ works differently than Kafka that works differently than WebHooks etc
  - Always dive deep into the docs of the channel you're using
  - We'll use RabbitMQ and SignalR in the implementation section

# Channel


- The Channel distributes the event to the Consumers



# Channel

---

- The channel's method of distribution varies between channels
- Can be:
  - Queue
  - REST API call
  - Proprietary listener



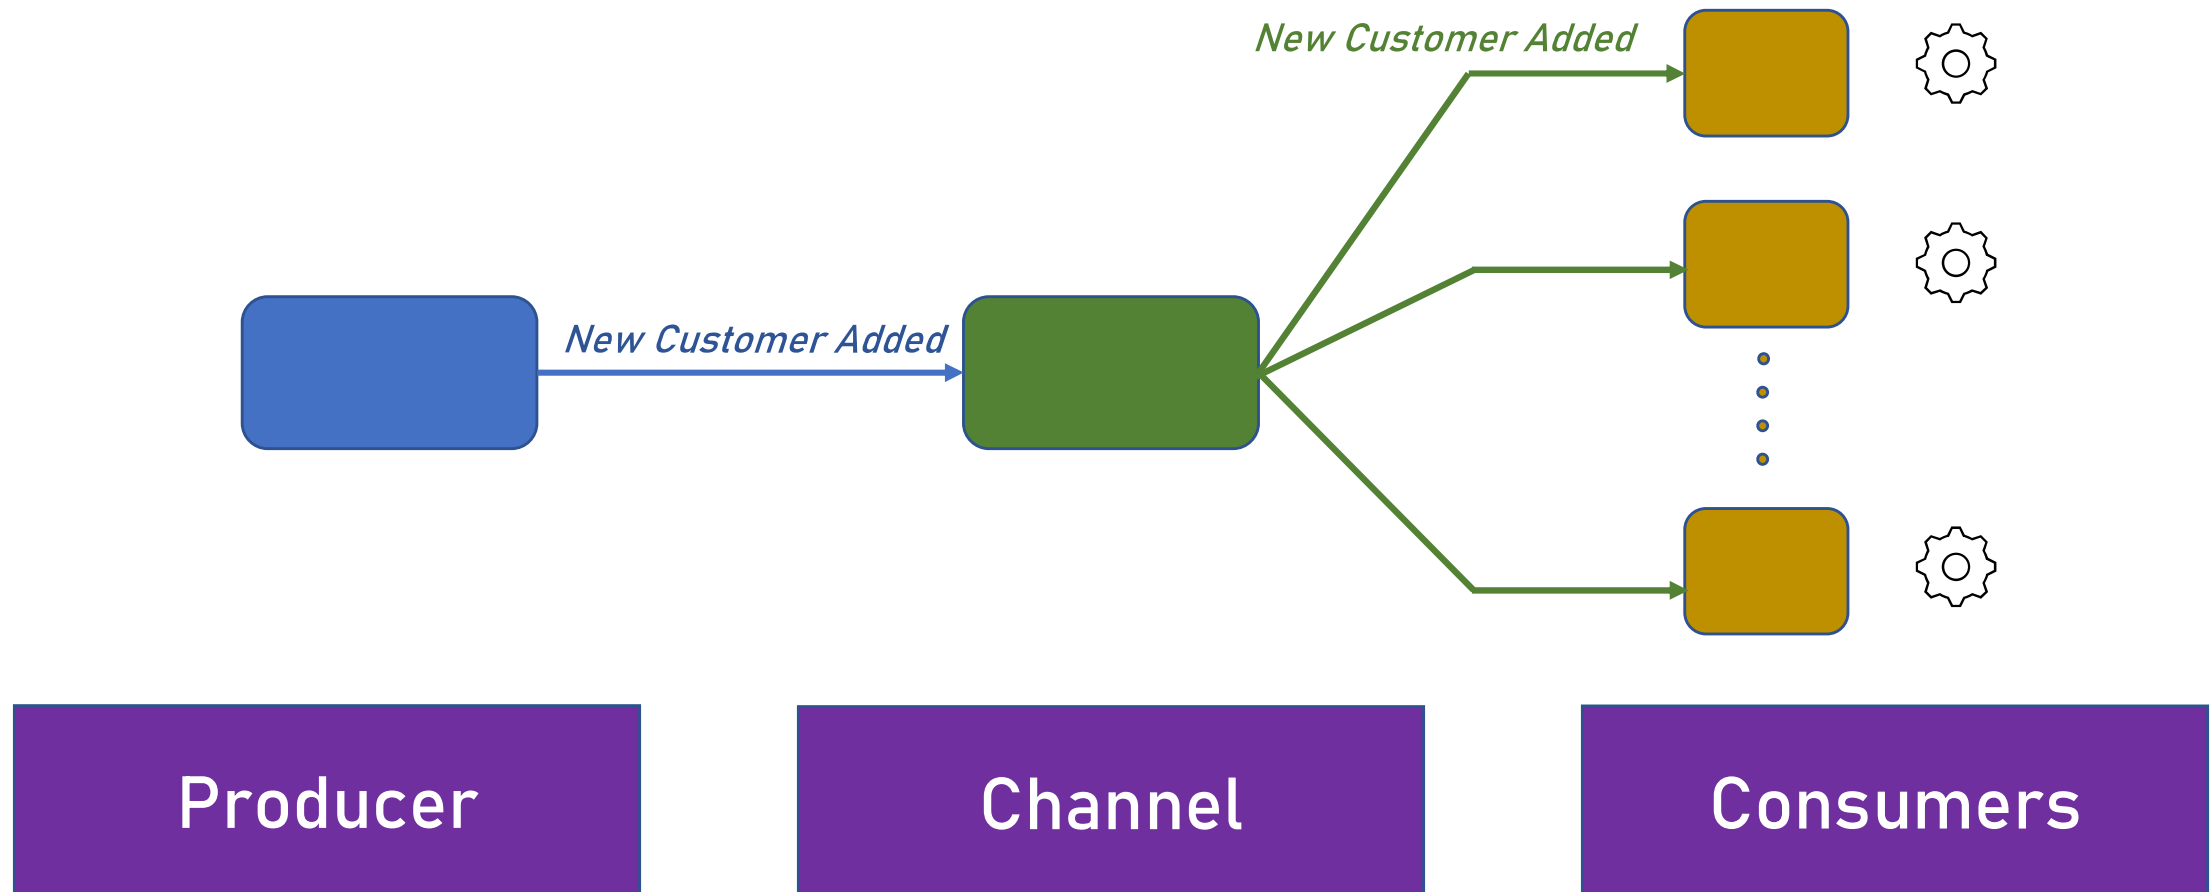
# Consumer

---

- The component that receives the event sent by the Producer and distributed by the Channel
- Can be developed in any development language compatible with the Channel's libraries (if any)
- Processes the event
- Sometimes – reports back when processing is complete (Ack)

# Consumer

- The Consumer receives and processes the event



# Consumer


---

- Consumer gets the event using either:
  - Push
  - Pull
- The method depends on the channel



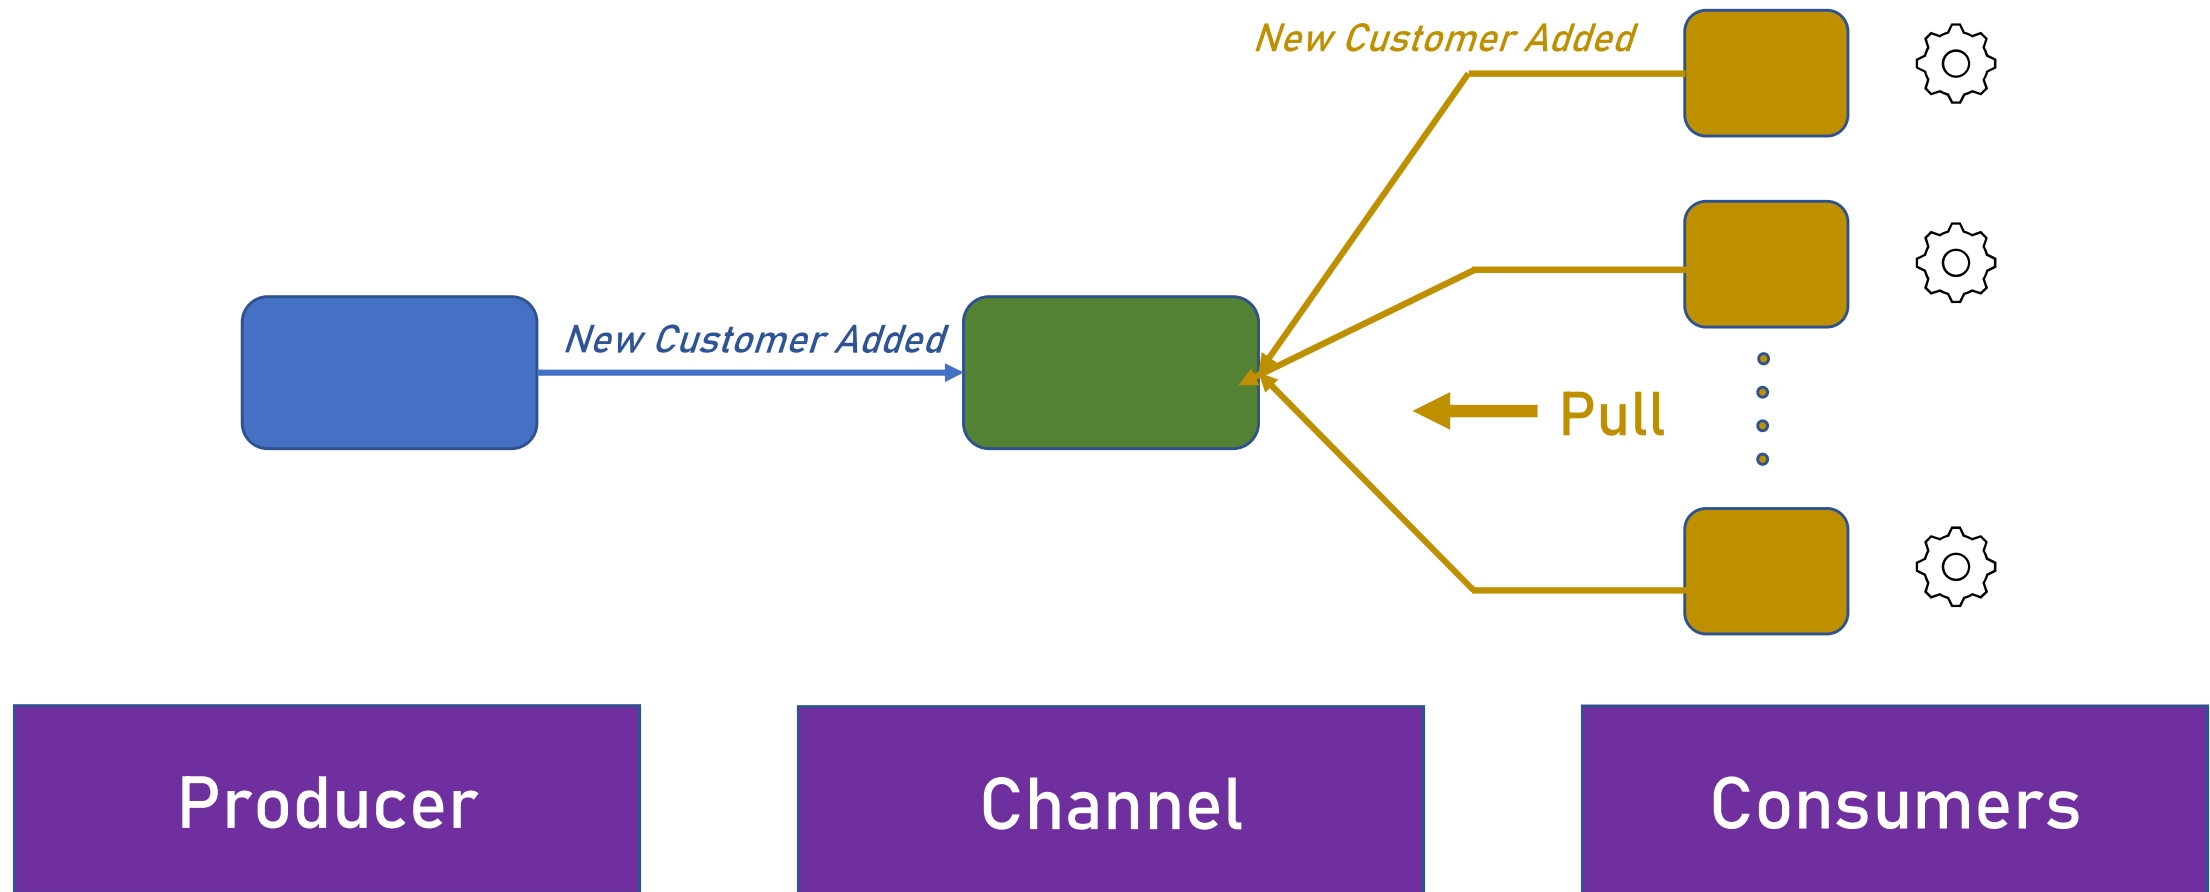
# Push

- The Channel pushes the event to the Consumers



# Pull

- The Consumers poll the Channel for new events



# Advantages of EDA

---

- Event Driven Architecture has a lot of advantages over other architecture paradigms
- As a quick refresher...

# Problems with Command and Query

---


- Three major problems with command and query:

Performance

Coupling

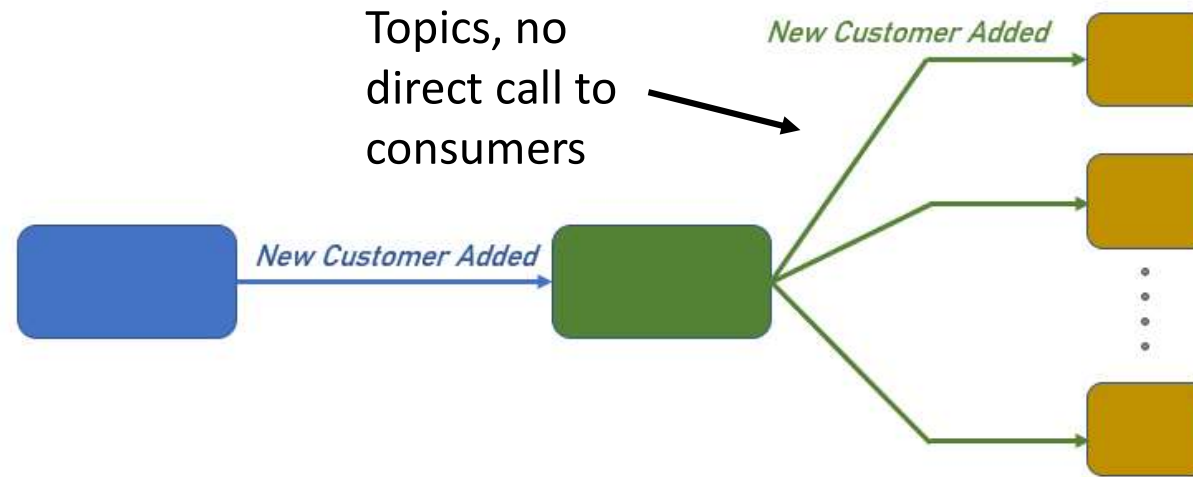
Scalability

# Performance




- EDA is an asynchronous architecture
- The Channel does not wait for response from consumer
- No performance bottlenecks

# Coupling



- The producer sends events to the channel
- The channel distributes events to topics / queues
- Both have no idea who's listening to the event (except in WebHooks)
- No coupling

# Scalability



- Many consumers can listen to events from channel
- More can be added as needed
- Channel doesn't care, producer doesn't know
- Fully scalable


# EDA and Pub/Sub

---

- Event Driven Architecture is often mentioned with Pub/Sub
- Pub/Sub = Publish and Subscribe
- A messaging pattern used by Event Driven Architecture



# Components of Pub/Sub



# EDA and Pub/Sub

---

- Event Driven Architecture and Pub/Sub are extremely similar
- Main difference:
  - EDA describes the whole architecture of the system
  - Pub/Sub is a messaging pattern used by the system
    - Not exclusively!


# EDA and Pub/Sub

---

- For example:
  - *“My Event Driven Architecture uses mainly Pub/Sub for inter-service communication, but I do have some REST APIs for synchronous queries.”*


# Ordering in EDA

- Messaging engines often guarantee the order of the messages
- Popular mainly in traditional queues



# Ordering in EDA

- With Event Driven Architecture (especially with Pub/Sub) ordering is not always guaranteed
- Ordering might be affected by consumer latency, code performance and more



# Ordering in EDA

---

- If ordering is important, make sure to select a channel that supports this capability
- Examples:
  - RabbitMQ supports it
  - SignalR does not
- We'll use both in the case study section

# Orchestration and Choreography

---

- Event Driven Architecture usually employs one of two architectural styles

Orchestration

Choreography


# Orchestration

---

- Flow of events in the system is determined by a central orchestrator
- Orchestrator receives output from components and calls the next component in the flow
- The next component sends the output back to the orchestrator etc.



# Orchestration



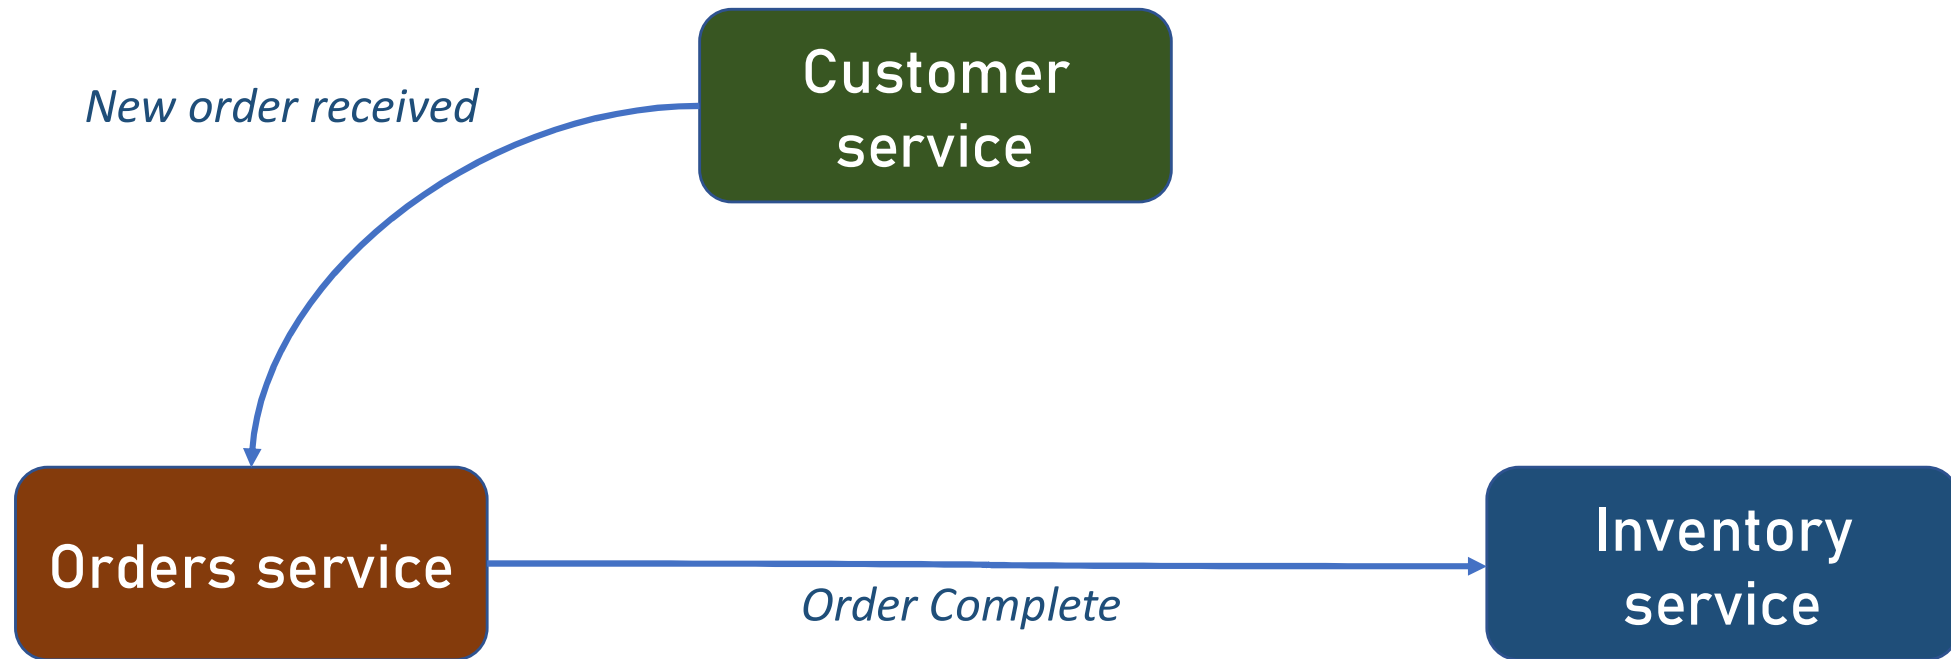
# Choreography

---

- No central “knowing all” component
- Each component notifies about the status of events
- Other components listen to the events and act accordingly

# Choreography

---



# Orchestration and Choreography

---

## Orchestration

- Logic is defined in a single place – easier to maintain
- Central traffic gateway – easier monitoring and logging

## Choreography

- Performance – no middleman
- Reliability – if one component fails, the rest still work

# Orchestration and Choreography

---

- Not constrained to EDA only
- Can be used with other types of communication
- Became popular with EDA

# Event Sourcing and CQRS

---

- Event Driven Architecture is mainly about services
- Events can be used as the basic building blocks of data too
- Event Sourcing and CQRS offer a pattern to store data as events

# Problems with Traditional DBs

---

- Traditional databases hold data about current state of entity
  - This is true for SQL and NoSQL databases

# Problems with Traditional DBs

- Example: Employees table

| emp_id | first_name | last_name | address                      | role                   | date_join  |
|--------|------------|-----------|------------------------------|------------------------|------------|
| 1      | John       | Smith     | Beverly Hills<br>90210       | Development<br>Manager | 2009-04-23 |
| 2      | Sarah      | Jones     | 42 <sup>nd</sup> st. NYC     | Sales                  | 2019-01-30 |
| 3      | Britney    | Flyn      | Marigold Lane,<br>Boca Raton | HR                     | 2022-05-19 |



# Problems with Traditional DBs

| emp_id | first_name | last_name | address                      | role                   | date_join  |
|--------|------------|-----------|------------------------------|------------------------|------------|
| 1      | John       | Smith     | Beverly Hills<br>90210       | Development<br>Manager | 2009-04-23 |
| 2      | Sarah      | Jones     | 42 <sup>nd</sup> st. NYC     | Sales                  | 2019-01-30 |
| 3      | Britney    | Flyn      | Marigold Lane,<br>Boca Raton | HR                     | 2022-05-19 |

- This table doesn't tell us:
  - What was John's previous role?
  - When did Sarah move to NYC?
  - Did any of the employees change his/her name?

# Problems with Traditional DBs

---

- Traditional databases hold data about current state of entity
- There is no way to see historical data of entities
- Data is a “snapshot” of a point in time
- Especially problematic with...

# Problems with Traditional DBs

| Date           | Description                       | Ref. | Withdrawals | Deposits | Balance |
|----------------|-----------------------------------|------|-------------|----------|---------|
| 2003-10-08     | Previous balance                  |      |             |          | 0.55    |
| 2003-10-14     | Payroll Deposit - HOTEL           |      |             | 694.81   | 695.36  |
| 2003-10-14     | Web Bill Payment - MASTERCARD     | 9685 | 200.00      |          | 495.36  |
| 2003-10-16     | ATM Withdrawal - INTERAC          | 3990 | 21.25       |          | 474.11  |
| 2003-10-16     | Fees - Interac                    |      | 1.50        |          | 472.61  |
| 2003-10-20     | Interac Purchase - ELECTRONICS    | 1975 | 2.99        |          | 469.62  |
| 2003-10-21     | Web Bill Payment - AMEX           | 3314 | 300.00      |          | 169.62  |
| 2003-10-22     | ATM Withdrawal - FIRST BANK       | 0064 | 100.00      |          | 69.62   |
| 2003-10-23     | Interac Purchase - SUPERMARKET    | 1559 | 29.08       |          | 40.54   |
| 2003-10-24     | Interac Refund - ELECTRONICS      | 1975 |             | 2.99     | 43.53   |
| 2003-10-27     | Telephone Bill Payment - VISA     | 2475 | 6.77        |          | 36.76   |
| 2003-10-28     | Payroll Deposit - HOTEL           |      |             | 694.81   | 731.57  |
| 2003-10-30     | Web Funds Transfer - From SAVINGS | 2620 |             | 50.00    | 781.57  |
| 2003-11-03     | Pre-Auth. Payment - INSURANCE     |      | 33.55       |          | 748.02  |
| 2003-11-03     | Cheque No. - 409                  |      | 100.00      |          | 648.02  |
| 2003-11-06     | Mortgage Payment                  |      | 710.49      |          | -62.47  |
| 2003-11-07     | Fees - Overdraft                  |      | 5.00        |          | -67.47  |
| 2003-11-08     | Fees - Monthly                    |      | 5.00        |          | -72.47  |
| *** Totals *** |                                   |      | 1,515.63    | 1,442.61 |         |

- Event Sourcing and CQRS try to solve this problem

# Event Sourcing

---

- A data store pattern in which every change in the data is captured and saved
- Database stores list of changes for the entity, not the entity itself
- No updates or deletes, just inserts
- Every row documents a change in a property/ies of the entity
- In this pattern, the database is called Event Store

# Event Sourcing

- Instead of this:

| emp_id | first_name | last_name | address                      | role                   | date_join  |
|--------|------------|-----------|------------------------------|------------------------|------------|
| 1      | John       | Smith     | Beverly Hills<br>90210       | Development<br>Manager | 2009-04-23 |
| 2      | Sarah      | Jones     | 42 <sup>nd</sup> st. NYC     | Sales                  | 2019-01-30 |
| 3      | Britney    | Flyn      | Marigold Lane,<br>Boca Raton | HR                     | 2022-05-19 |

# Event Sourcing

- We have this:

| event_id | timestamp  | event   |
|----------|------------|---|
| 1        | 2009-04-23 | Employee John Smith joined                            |
| 2        | 2009-04-23 | Address of John Smith updated to Hott Street, Clinton |
| 3        | 2009-04-23 | Role of John Smith updated to Junior Developer        |
| 4        | 2013-05-22 | Address of John Smith updated to Beverly Hills 90210  |
| 5        | 2017-09-12 | Role of John Smith updated to Development Manager     |
| 6        | 2019-01-30 | Employee Sarah Jones joined                           |
| 7        | 2019-01-30 | Role of Sarah Jones updated to Sales                  |
| 8        | 2021-07-05 | David Richer left the company                         |

- Specific columns are up to you, depends on the system requirements
- Note there's a lot more information than in the regular table

# Event Sourcing

- How can we view the current state of an entity?
- By replaying the events

| event_id | timestamp  | event   |
|----------|------------|---|
| 1        | 2009-04-23 | Employee John Smith joined                            |
| 2        | 2009-04-23 | Address of John Smith updated to Hott Street, Clinton |
| 3        | 2009-04-23 | Address of John Smith updated to Beverly Hills        |
| 4        | 2009-04-23 | John Smith became a Development Manager               |
| 5        | 2017-09-12 | Role of John Smith updated to Development Manager     |
| 6        | 2019-01-30 | Employee Sarah Jones joined                           |
| 7        | 2019-01-30 | Role of Sarah Jones updated to Sales                  |
| 8        | 2021-07-05 | David Richer left the company                         |

| emp_id | first_name | last_name | address             | role                | date_join  |
|--------|------------|-----------|---------------------|---------------------|------------|
| 1      | John       | Smith     | Beverly Hills 90210 | Development Manager | 2009-04-23 |

# Event Sourcing

## Pros

- Extremely easy to view historical data
- Simple database structure
- Simple database operations (no updates, no concurrency)
- Very fast inserts

## Cons

- Viewing current entity state is cumbersome and slow
- Large database capacity (many records per entity)

CQRS to the Rescue!



# CQRS

---

- Stands for:

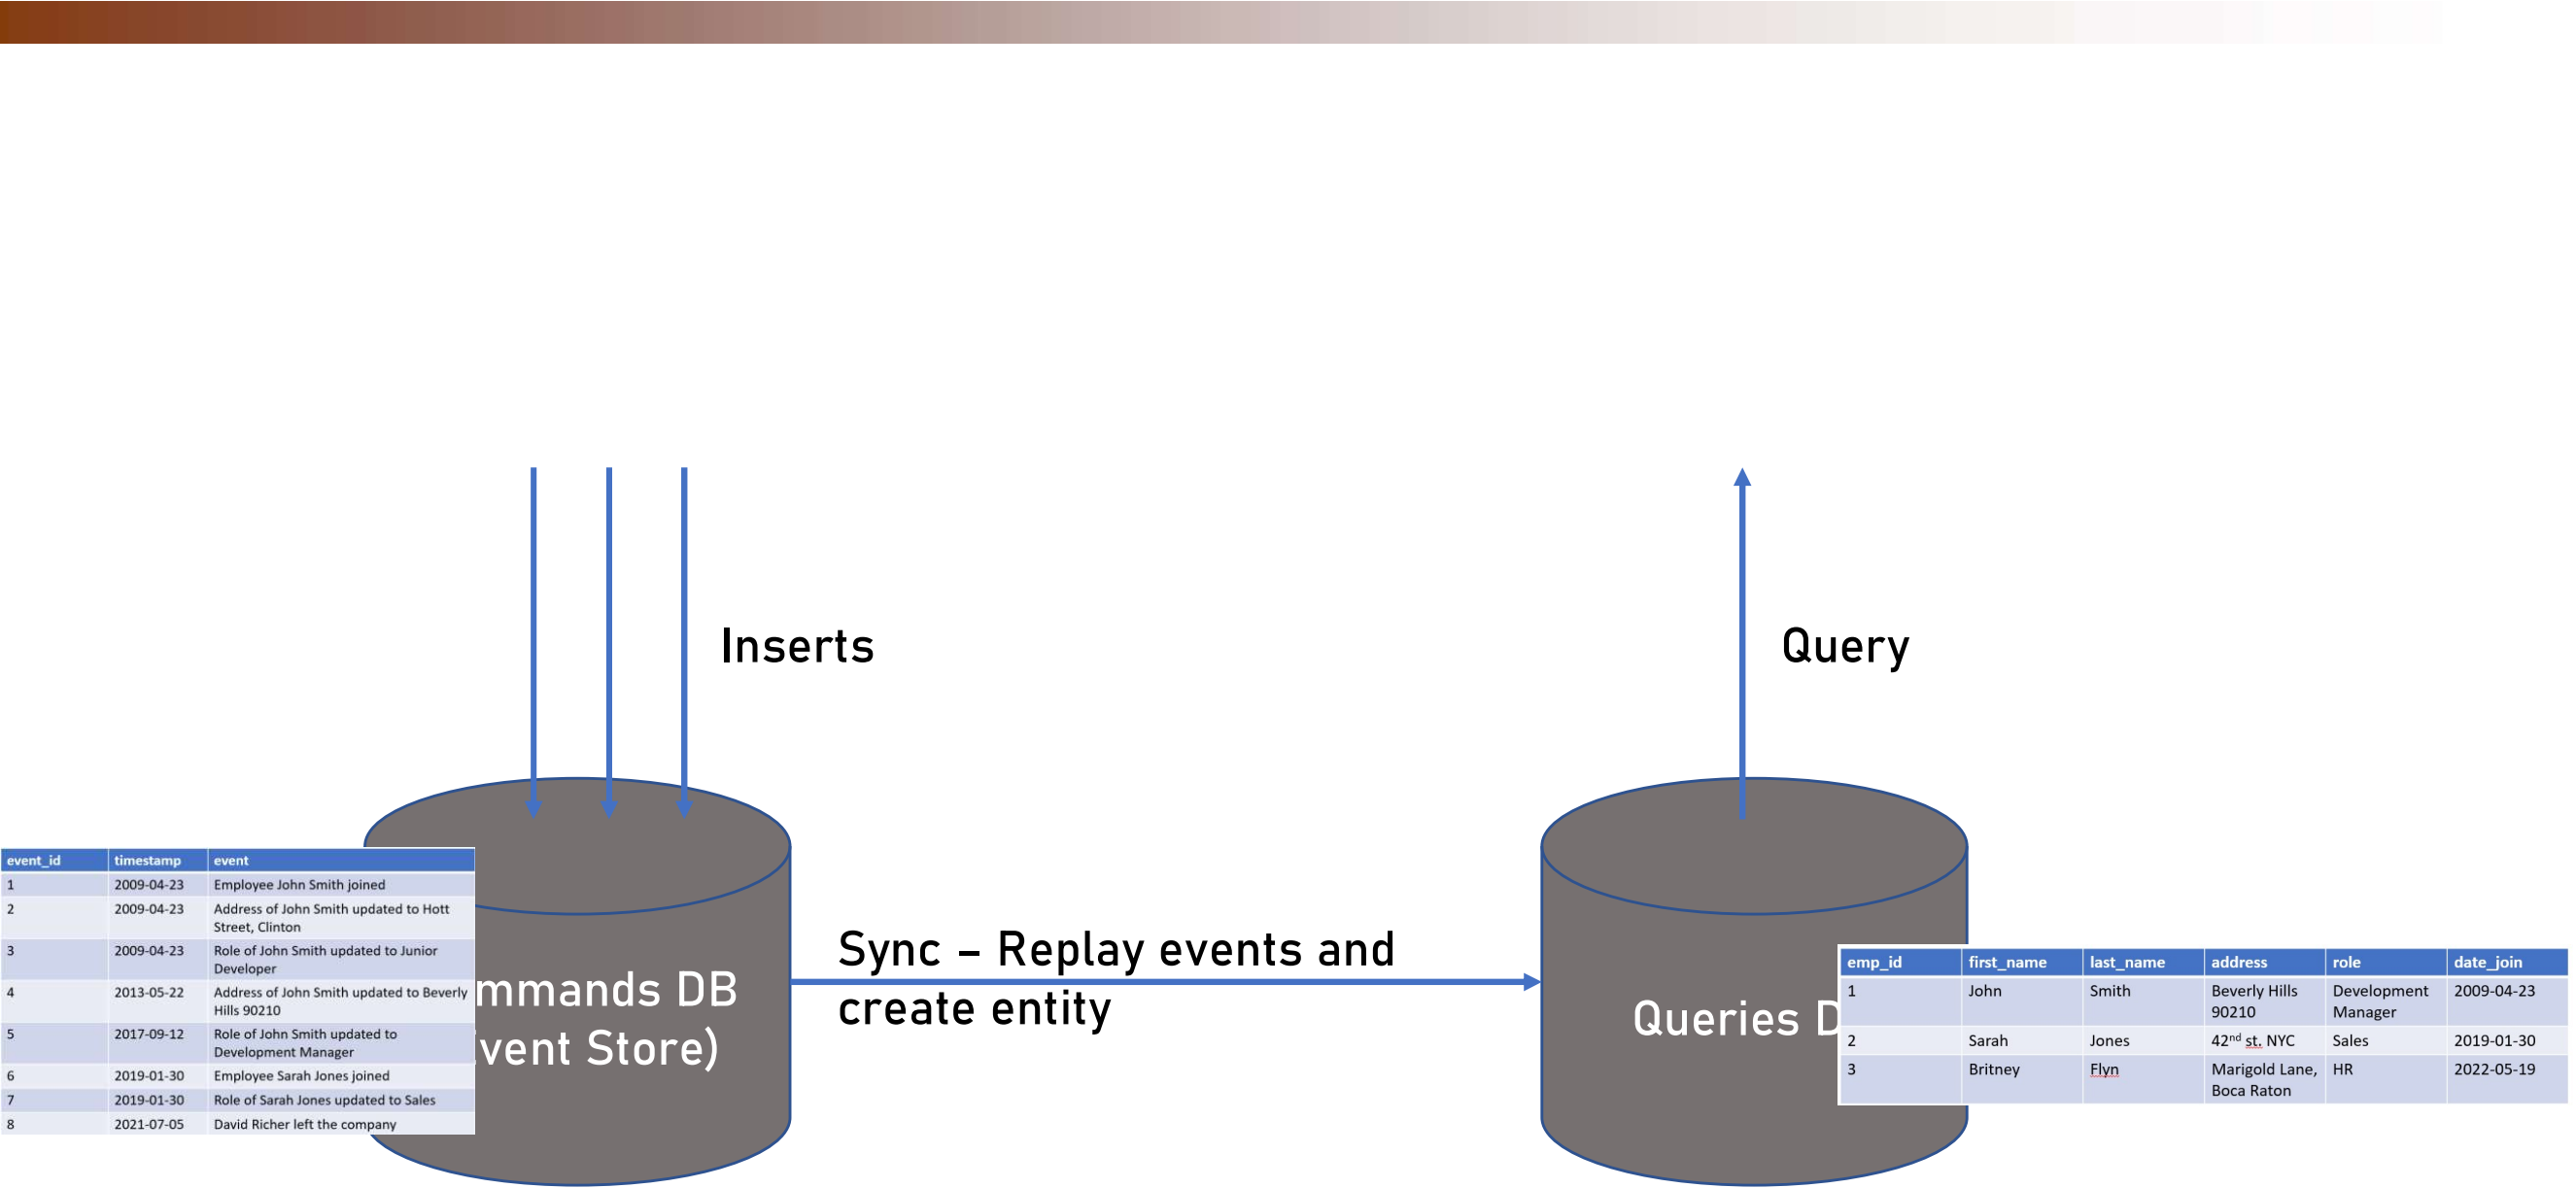
Command and Query Responsibility Segregation

# CQRS

---

- **Means:**
  - Separating the commands (updates / inserts / deletes) from the queries
  - Each one of them in a separate database
  - Commands database is implemented as Event Store to improve performance and simplicity
  - Queries database stores entities
  - Database are synced using a central synchronization mechanism

# CQRS



# CQRS

---

## Pros

- Combines Event Sourcing pros with traditional entity query
- No performance hit when querying entities

## Cons

- Entity data is not updated in real-time
- Difficult to set-up and maintain

# When to Use Event Sourcing & CQRS

---

- When access to historical data is extremely important
  - Regulation, finance, healthcare etc.
- When data is large and replaying events is not feasible
- When performance is critical (inserts or queries)

# When to Use Event Driven Architecture

---

- Event Driven Architecture is not easy to implement
- Requires setting up and configuring channels
- Not trivial logging and monitoring
- Be sure to use it when needed


# Scalability

---

- Scalability is a non-issue in EDA
- New consumers can be added as needed with no changes to the architecture
- Great for fluctuating load

# Scalability

---





# Asynchronous

---

- If inter-service communication can be asynchronous, consider EDA
- Remember: EDA is async by nature
- Examples:
  - Send instructions to perform payment
  - Write to log


# Asynchronous

---

- Check how many synchronous interactions there are
- Usually mainly queries
- The more synchronous calls – the less EDA is relevant

# Reliable Network

- EDA utilizes a lot of traffic



- Network should be reliable or performance will be slow

# When not to Use EDA

---

- EDA is not suitable for:
  - Small systems with a few services
  - Synchronous-oriented systems
    - ie. Information system serving mainly queries from end users

# Stateless vs Stateful EDA

---

- There are two main patterns in implementing EDA
- Stateless and Stateful
- Related to the consumers behavior
- Both are legitimate, but make sure to select the right one for your scenario

# Stateless vs Stateful EDA

---

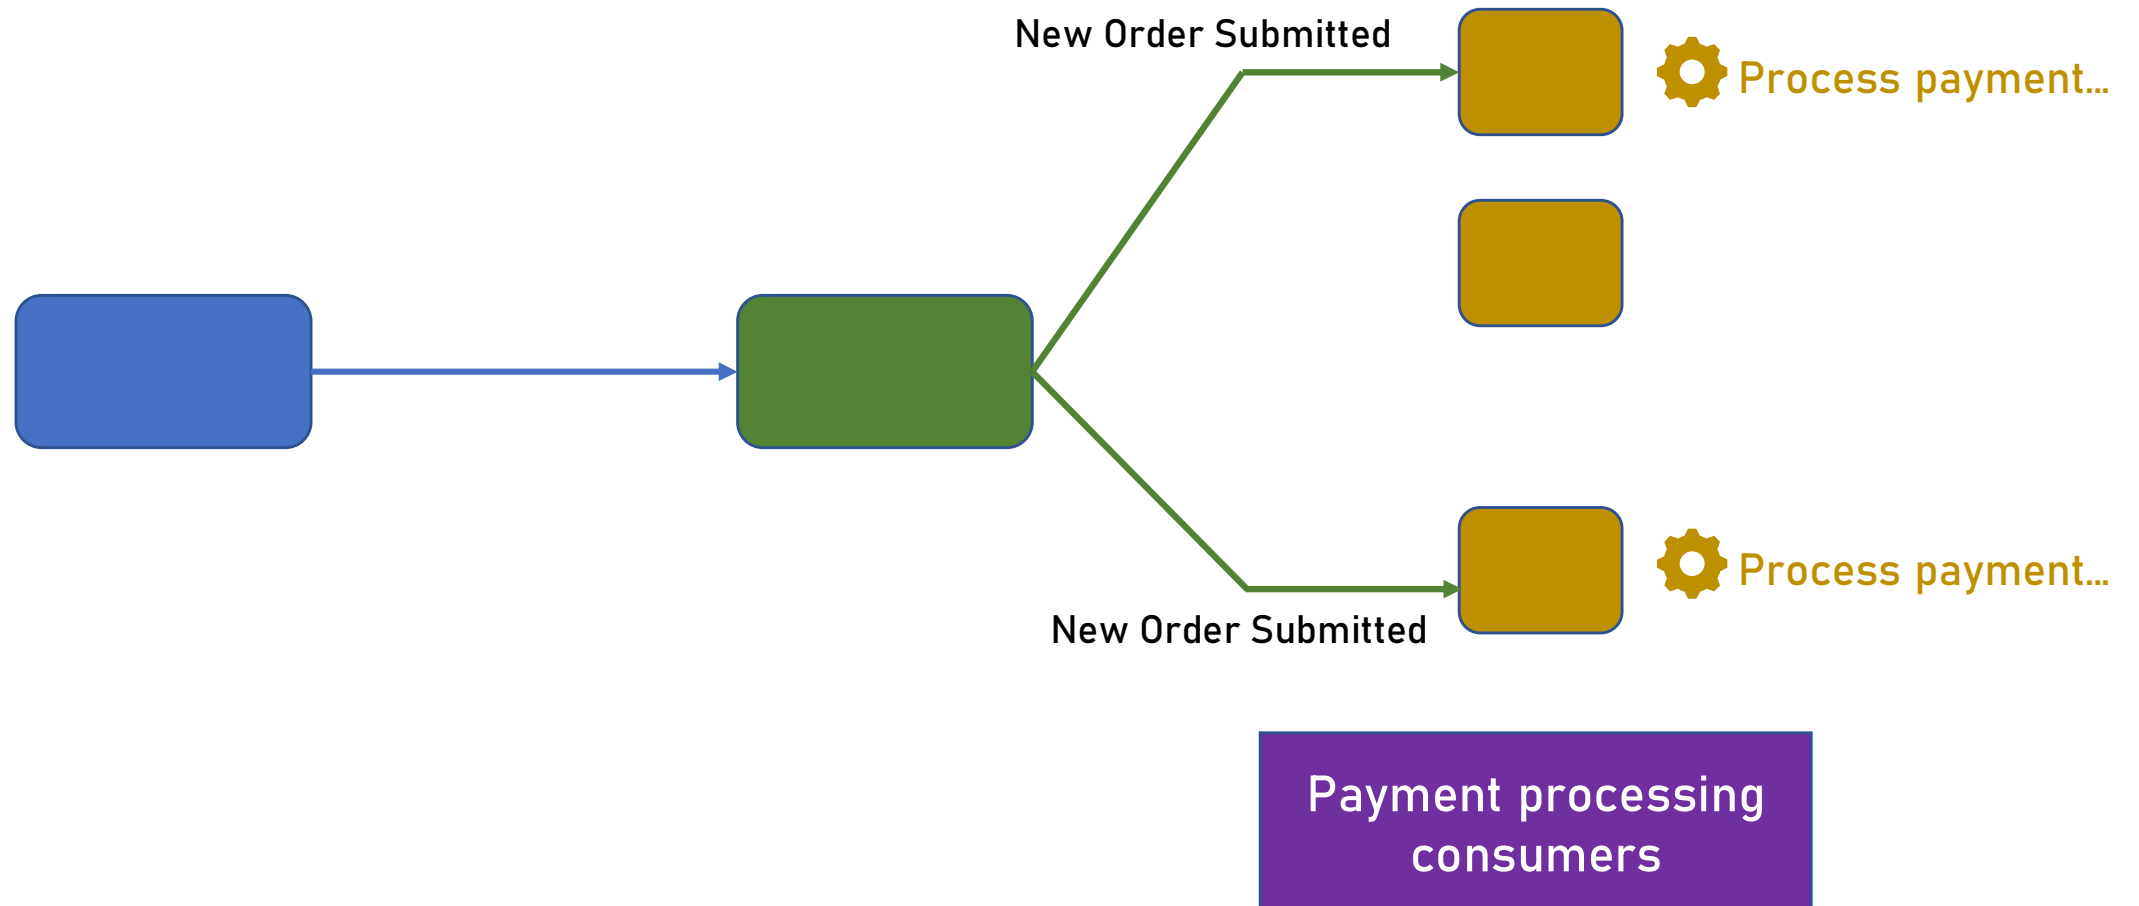
- In software architecture there's also the stateless vs stateful debate
- While the concepts are similar, the reasoning is different
- With software architecture it's often said that:
  - “Stateful is bad”
- This is not necessarily the case with EDA

# Stateless EDA

---

- Each event handled by a consumer is completely autonomous and is not related to past / future events
- Should be used when the event is an independent unit with its own outcomes

# Stateless EDA





# Stateless EDA

---

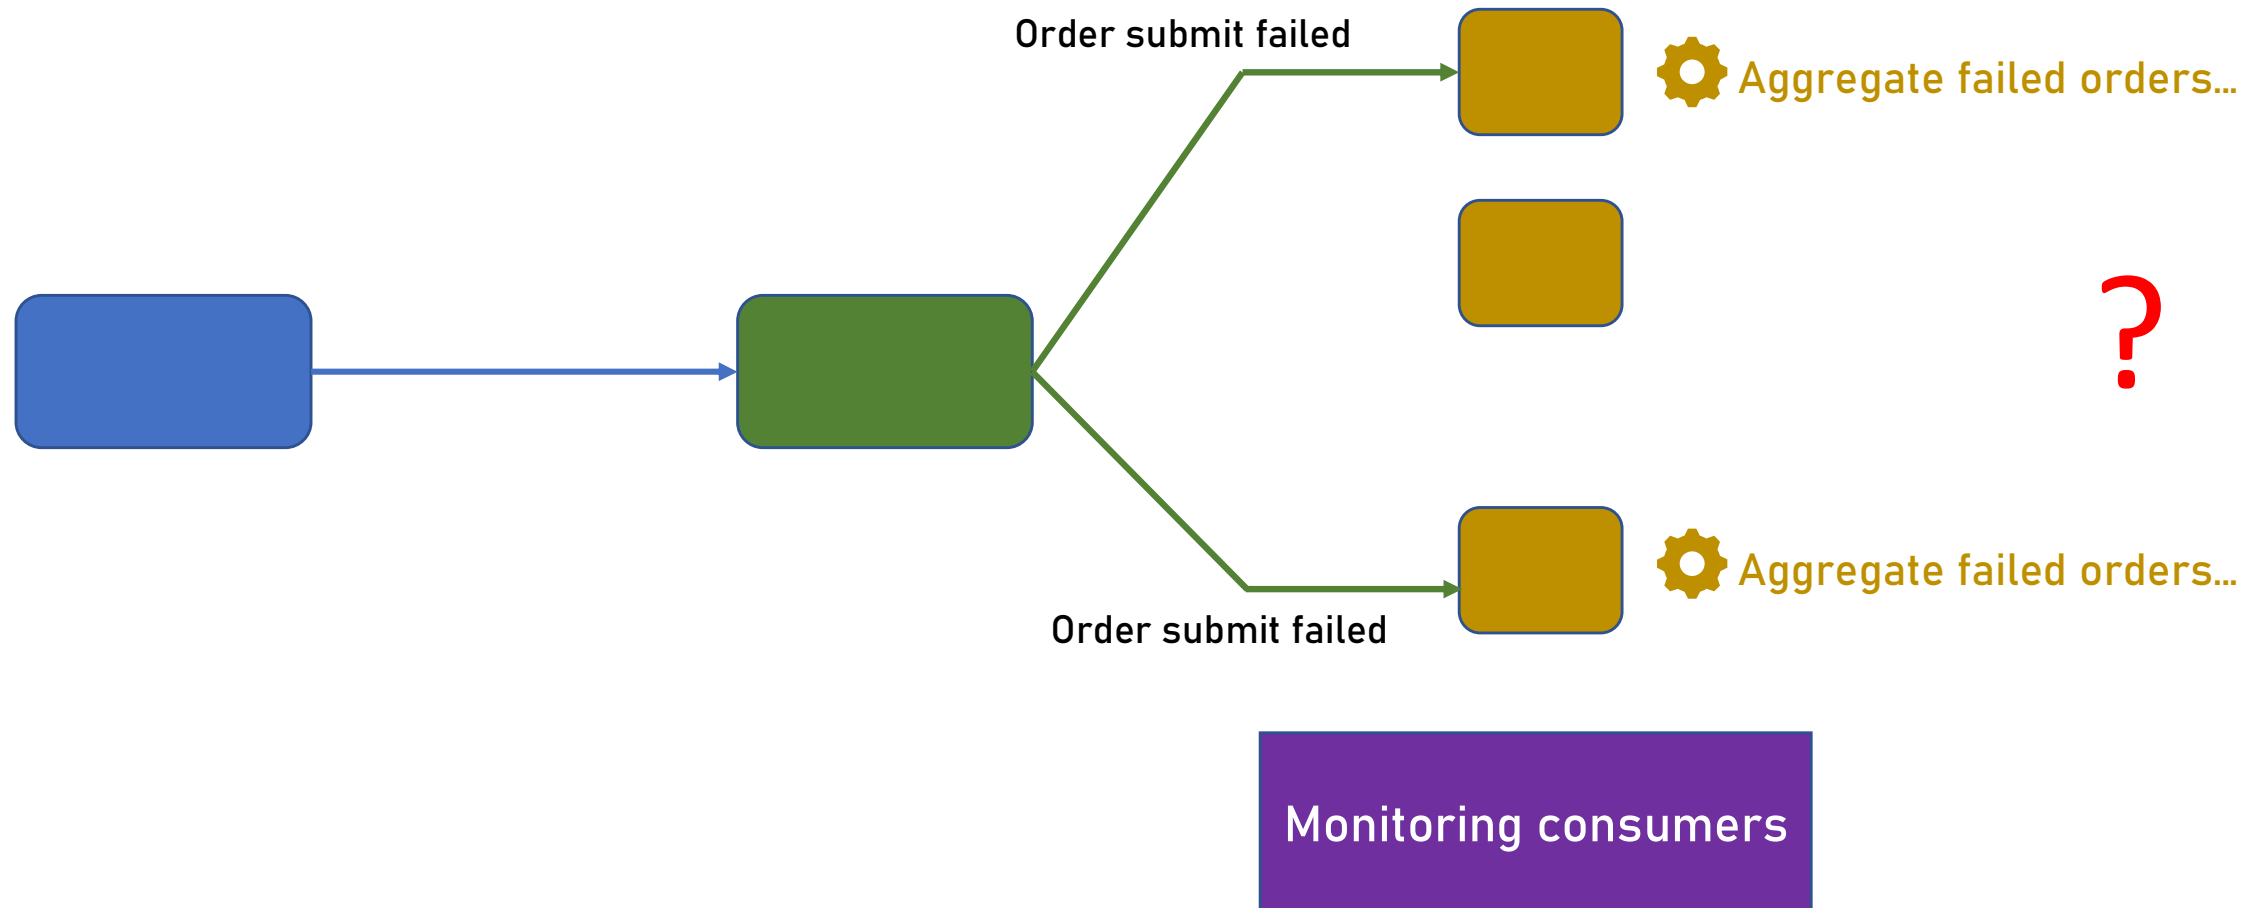
- It doesn't matter which consumer is handling the event
- The outcome is always the same
- Should be used when each event is autonomous
- Note: Has nothing to do with the question of what data is contained in the event and whether a call to a DB is required

# Stateful EDA

---

- Events might be related to past / future events
- Should be used mainly for aggregators and time-related events
- Examples:
  - Send an email if more than 5 failure events were received in a single minute
  - Calculate the amount of orders submitted in an hour

# Stateful EDA



# Stateful EDA

---

- It's extremely important which consumer handles the event
- Current state is stored in specific consumer(s)
- Should be used when events are part of a chain of events

# Problems with Stateful EDA

---

- Stateful EDA presents some problems that should be taken care of

Load balancing

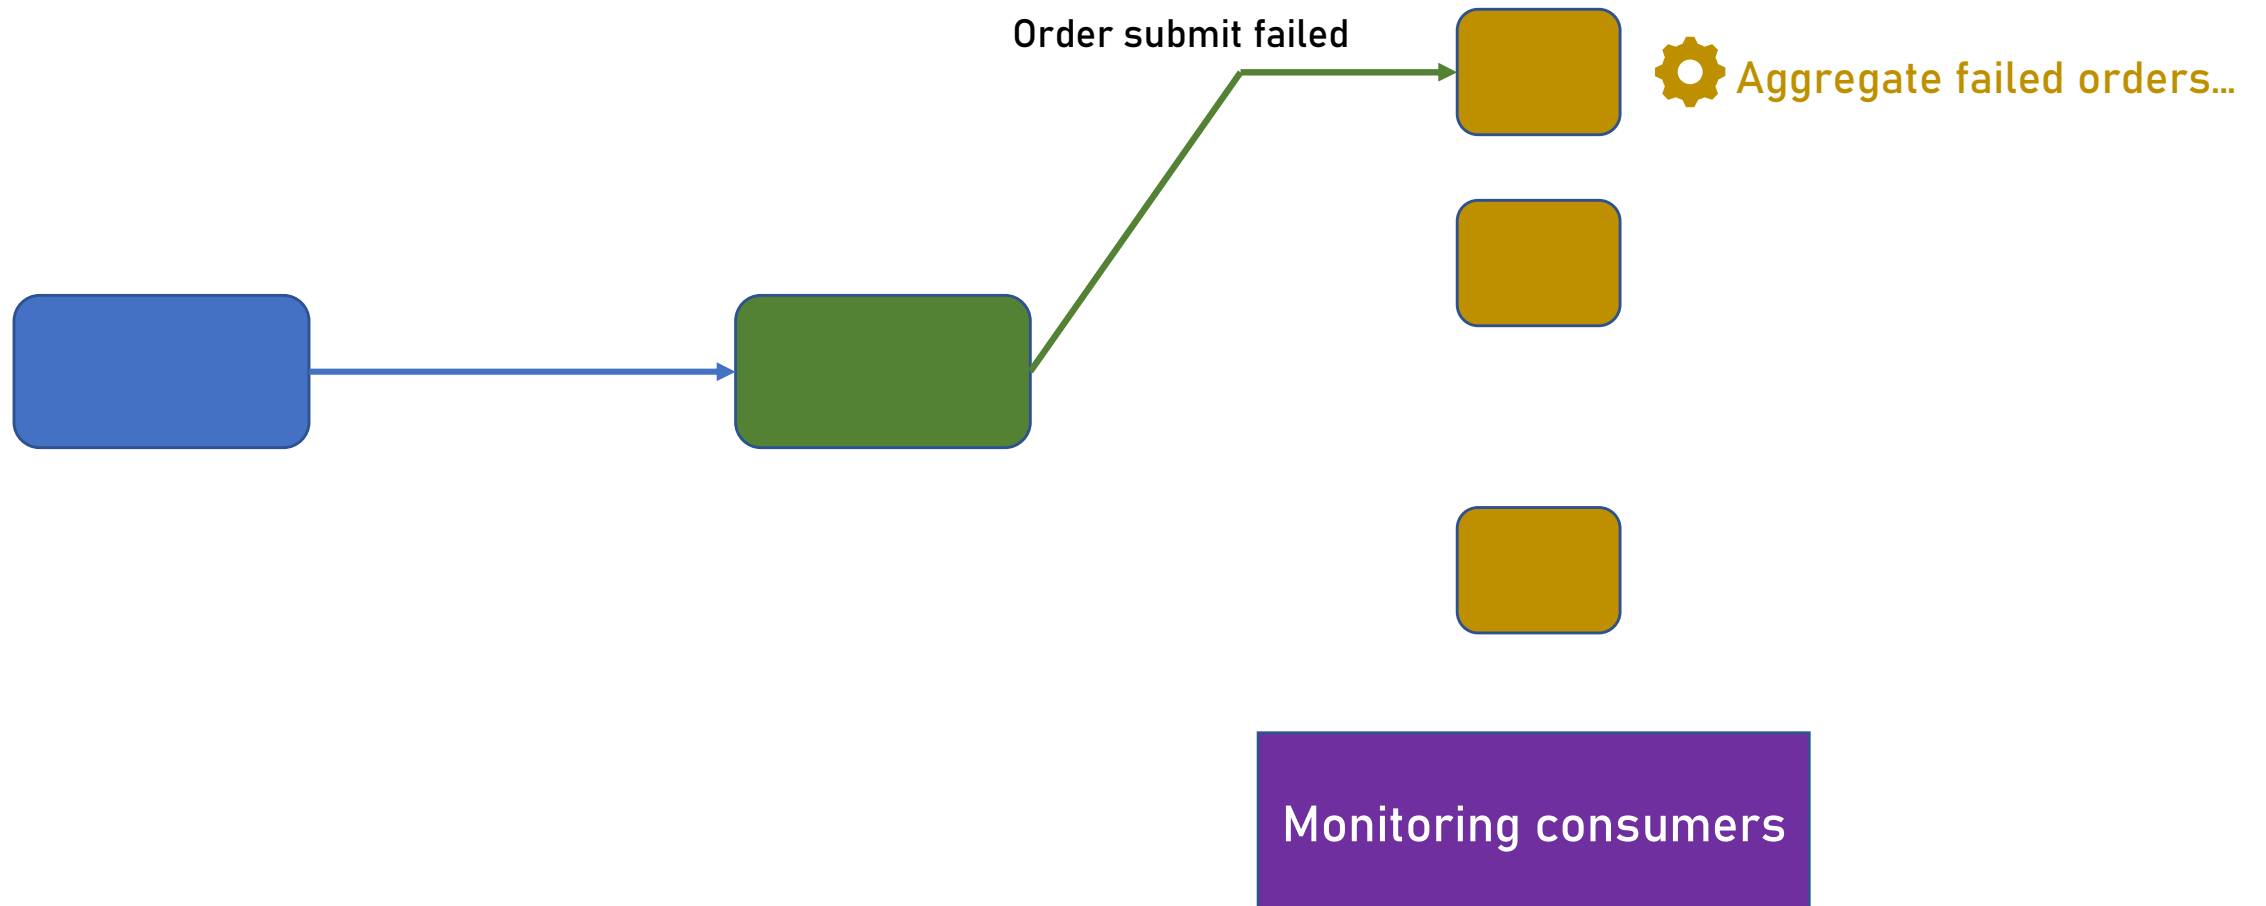
Scalability

# Load Balancing


---

- Since the state is stored in a specific consumer, subsequent events must be routed to the same consumer
- No load balancing is possible

# Load Balancing

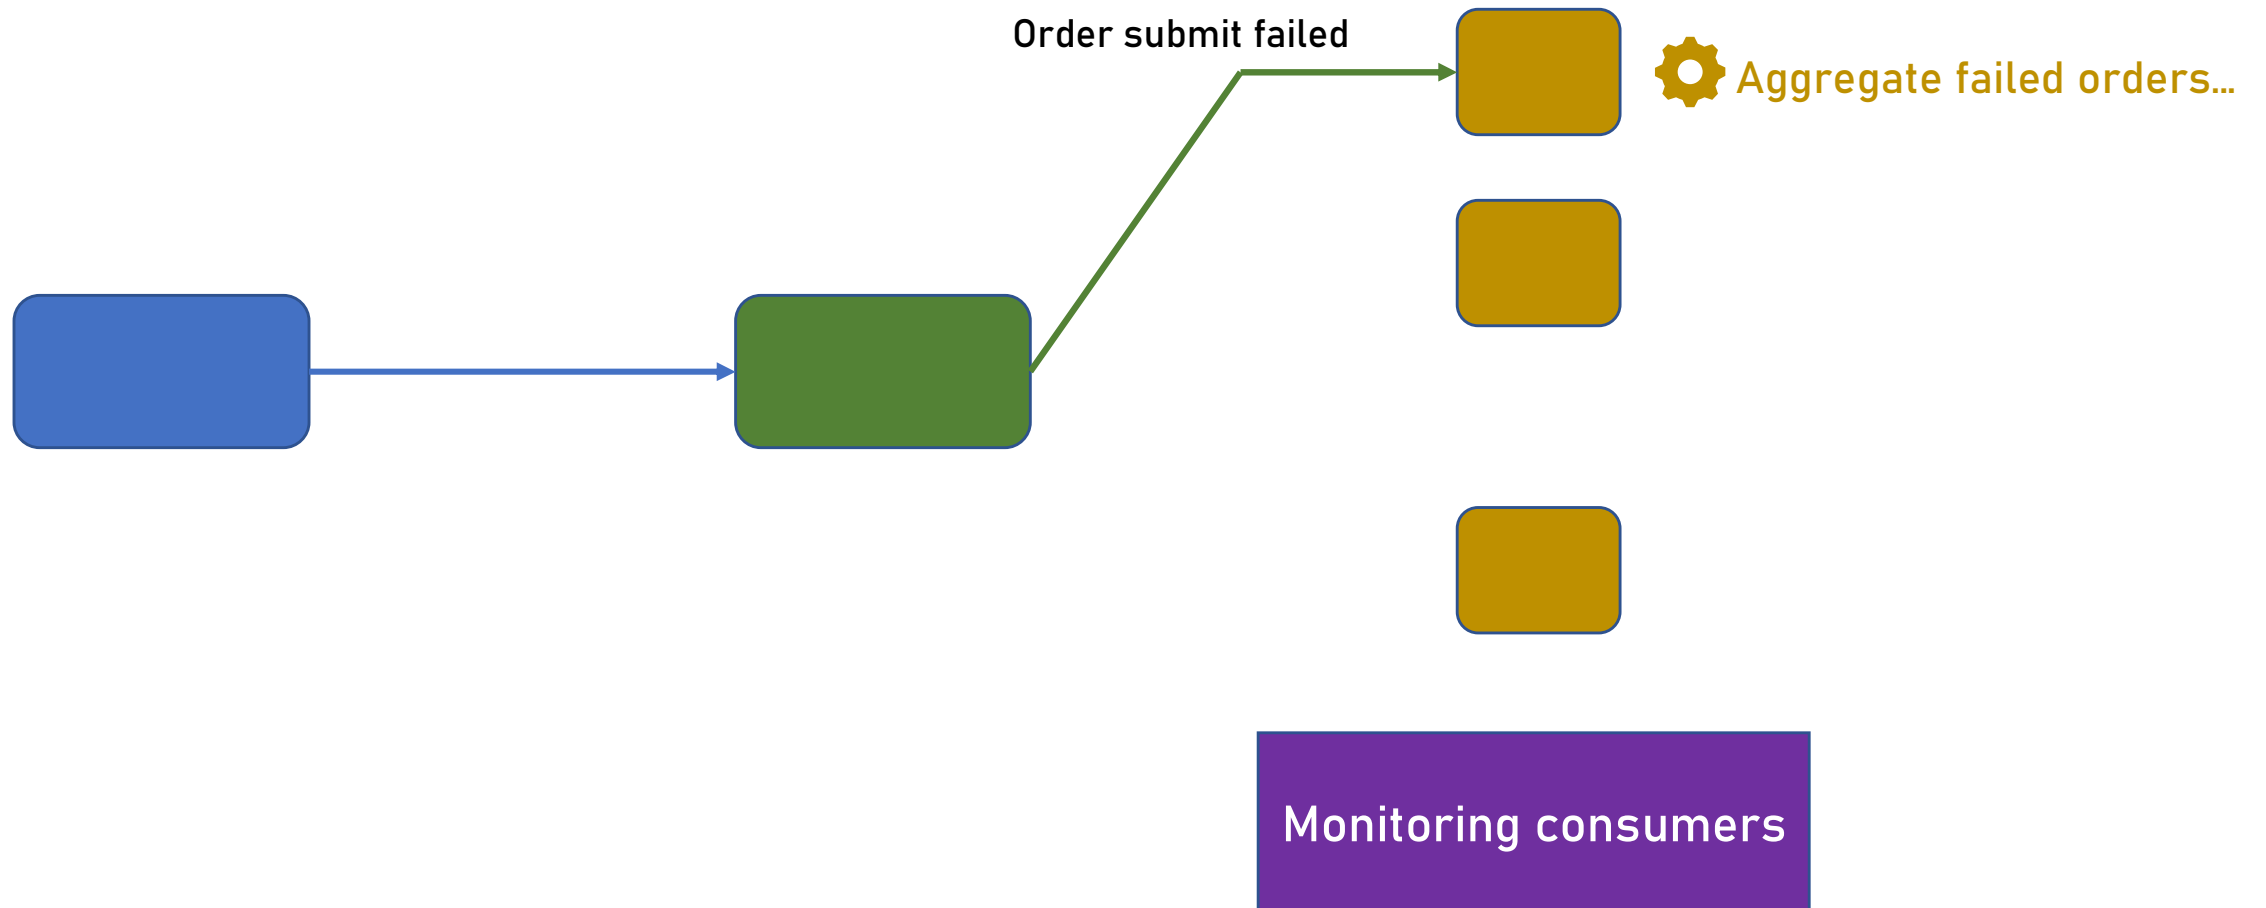


# Load Balancing





# Load Balancing

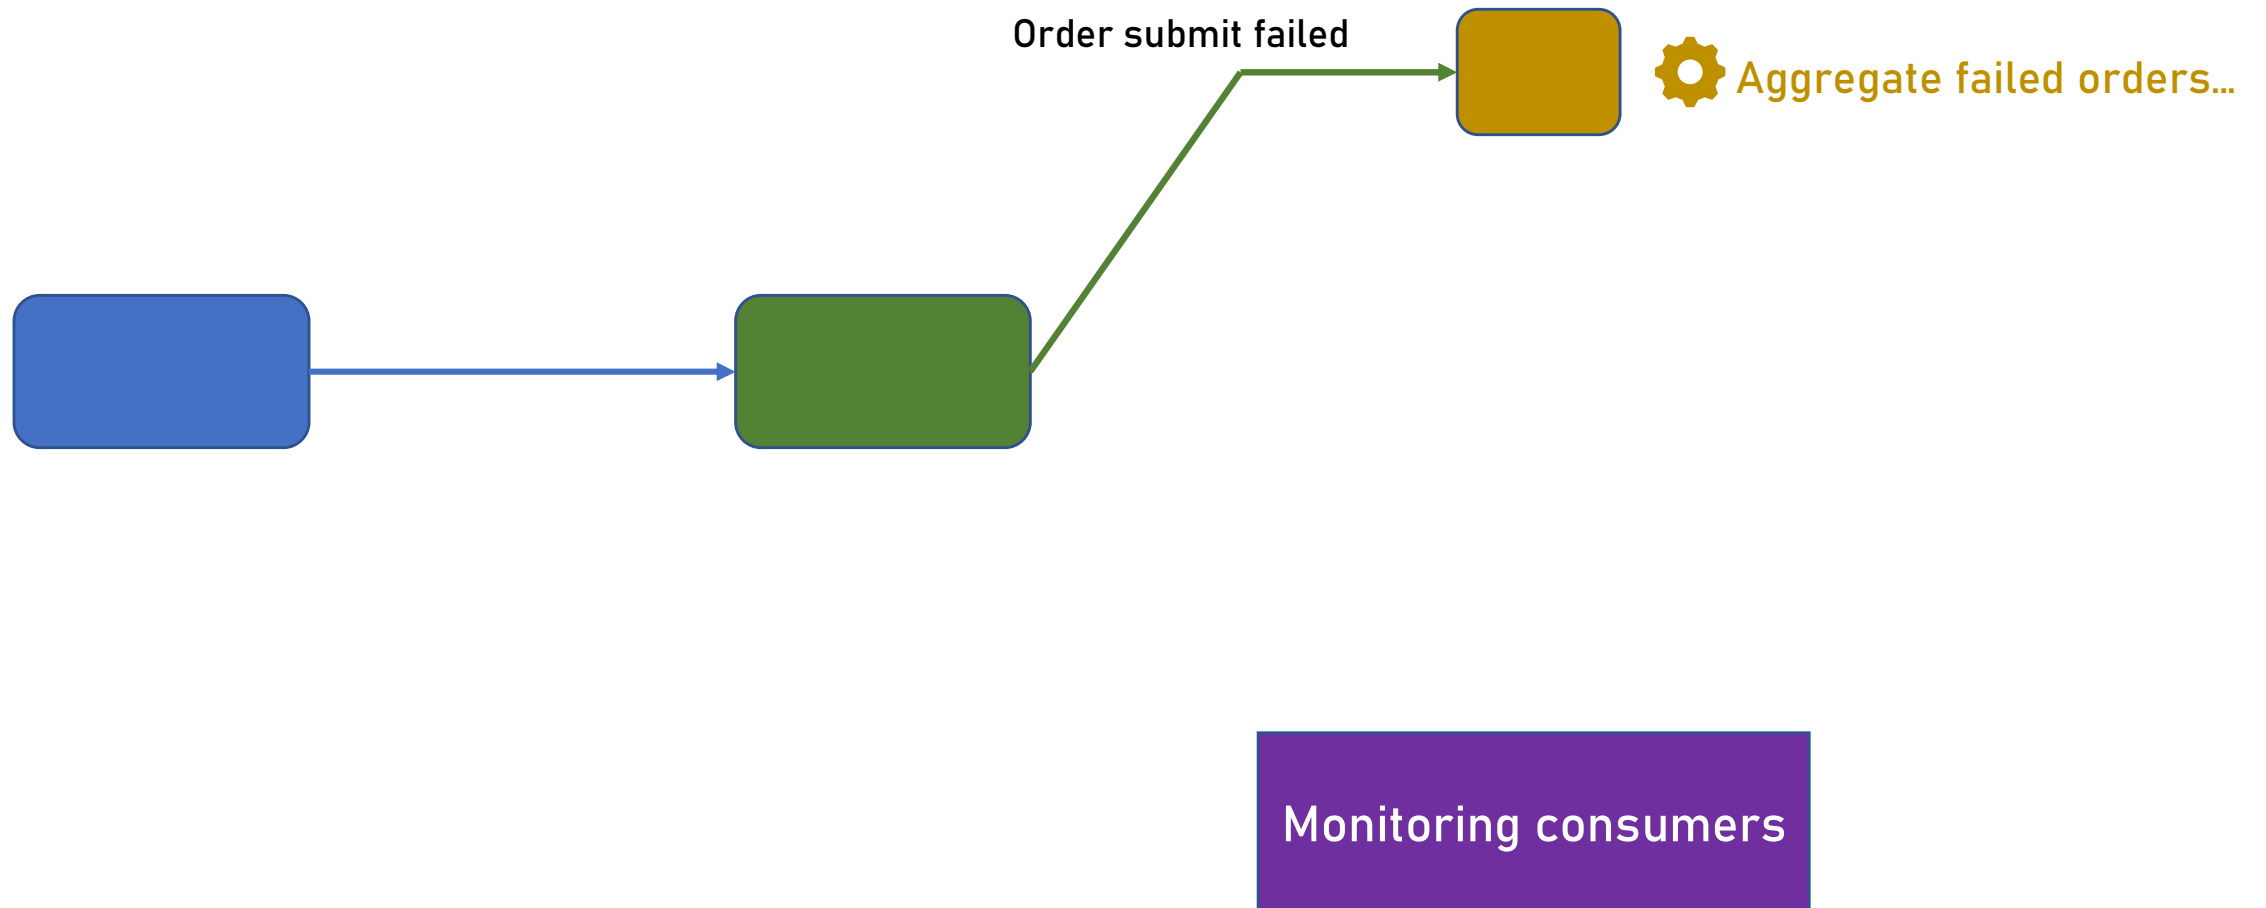


# Scalability

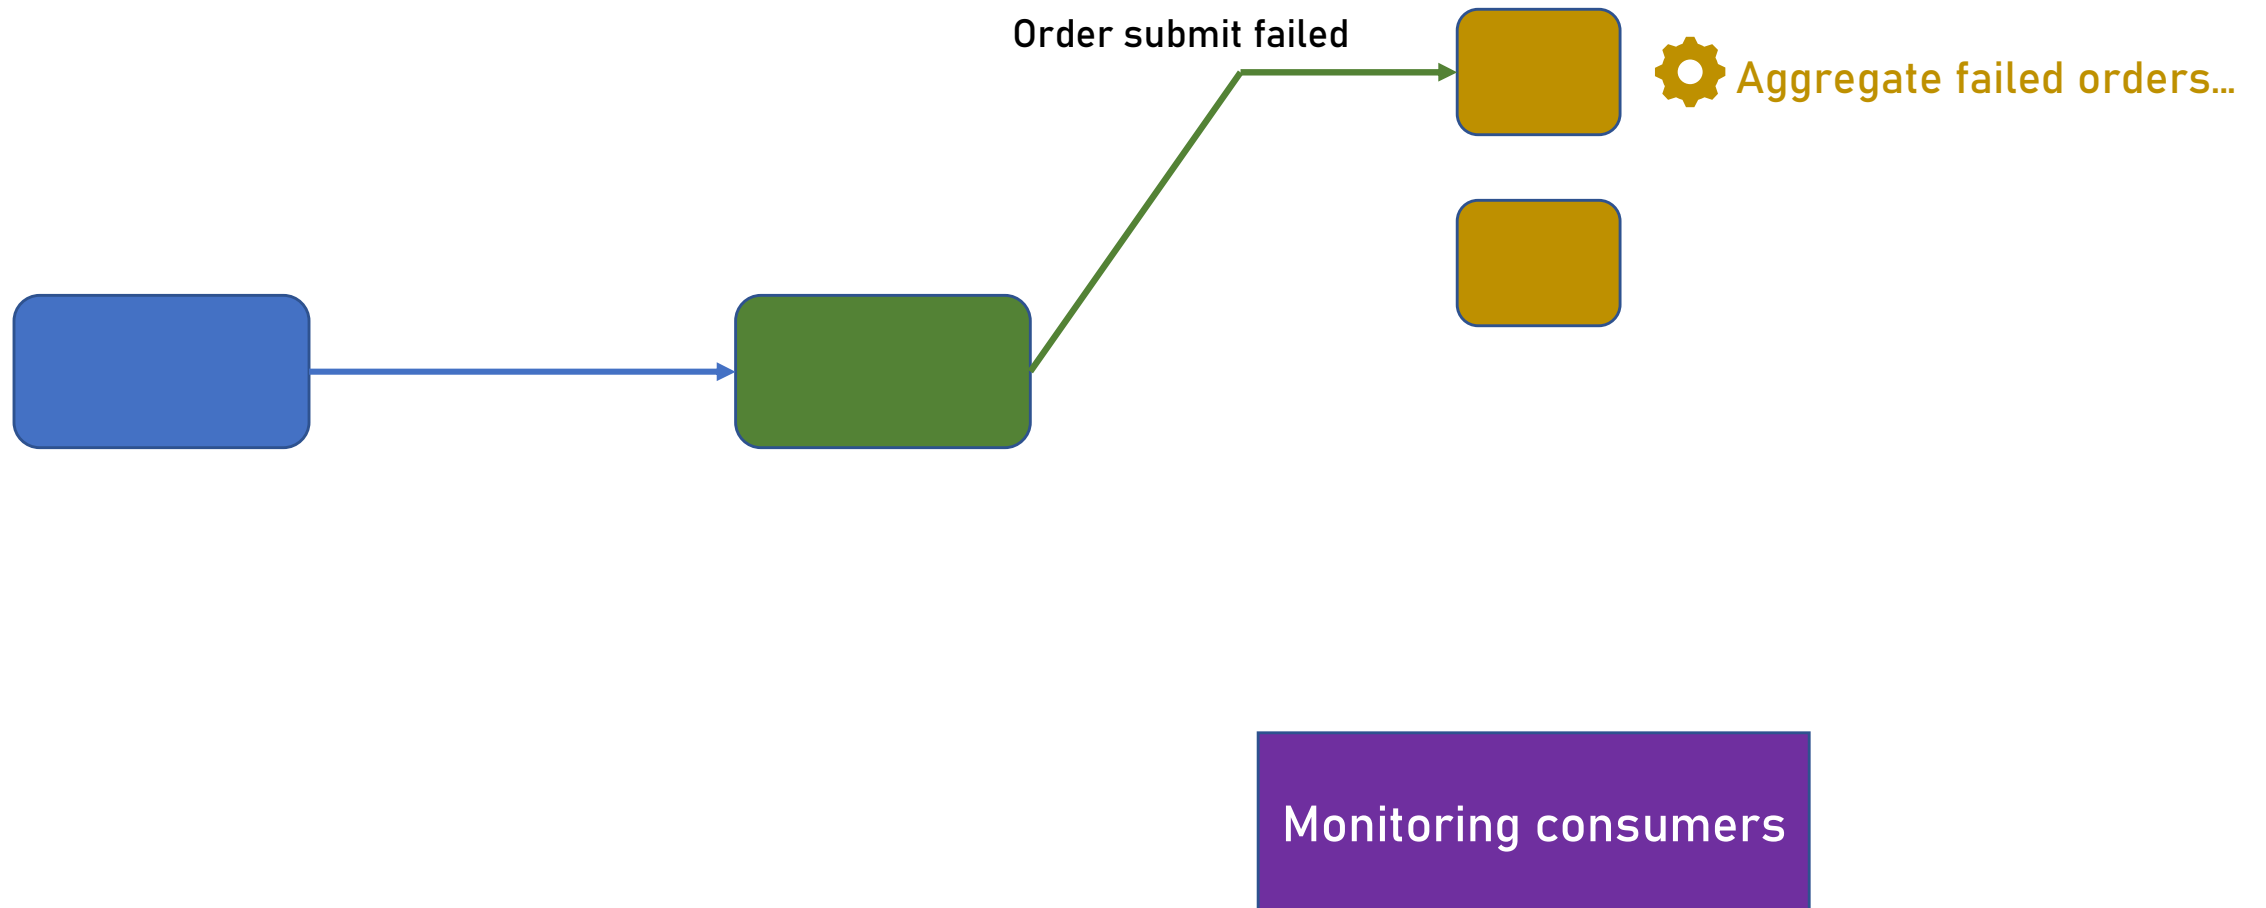
---

- Since the state is stored in a specific consumer, additional consumers cannot be added to handle the events

# Scalability



# Scalability



# Stateless vs Stateful

---

- Rule of thumb:
  - Use stateless EDA unless the business requirements force you to use stateful

# Event Streaming

---


- So far we talked about Event Driven Architecture
  - Something happened
  - An event was created
  - Someone listened to the event and handled it
- Event Streaming is another event-oriented pattern
- These are not the same, but share similar characteristics

# What is Event Streaming

---

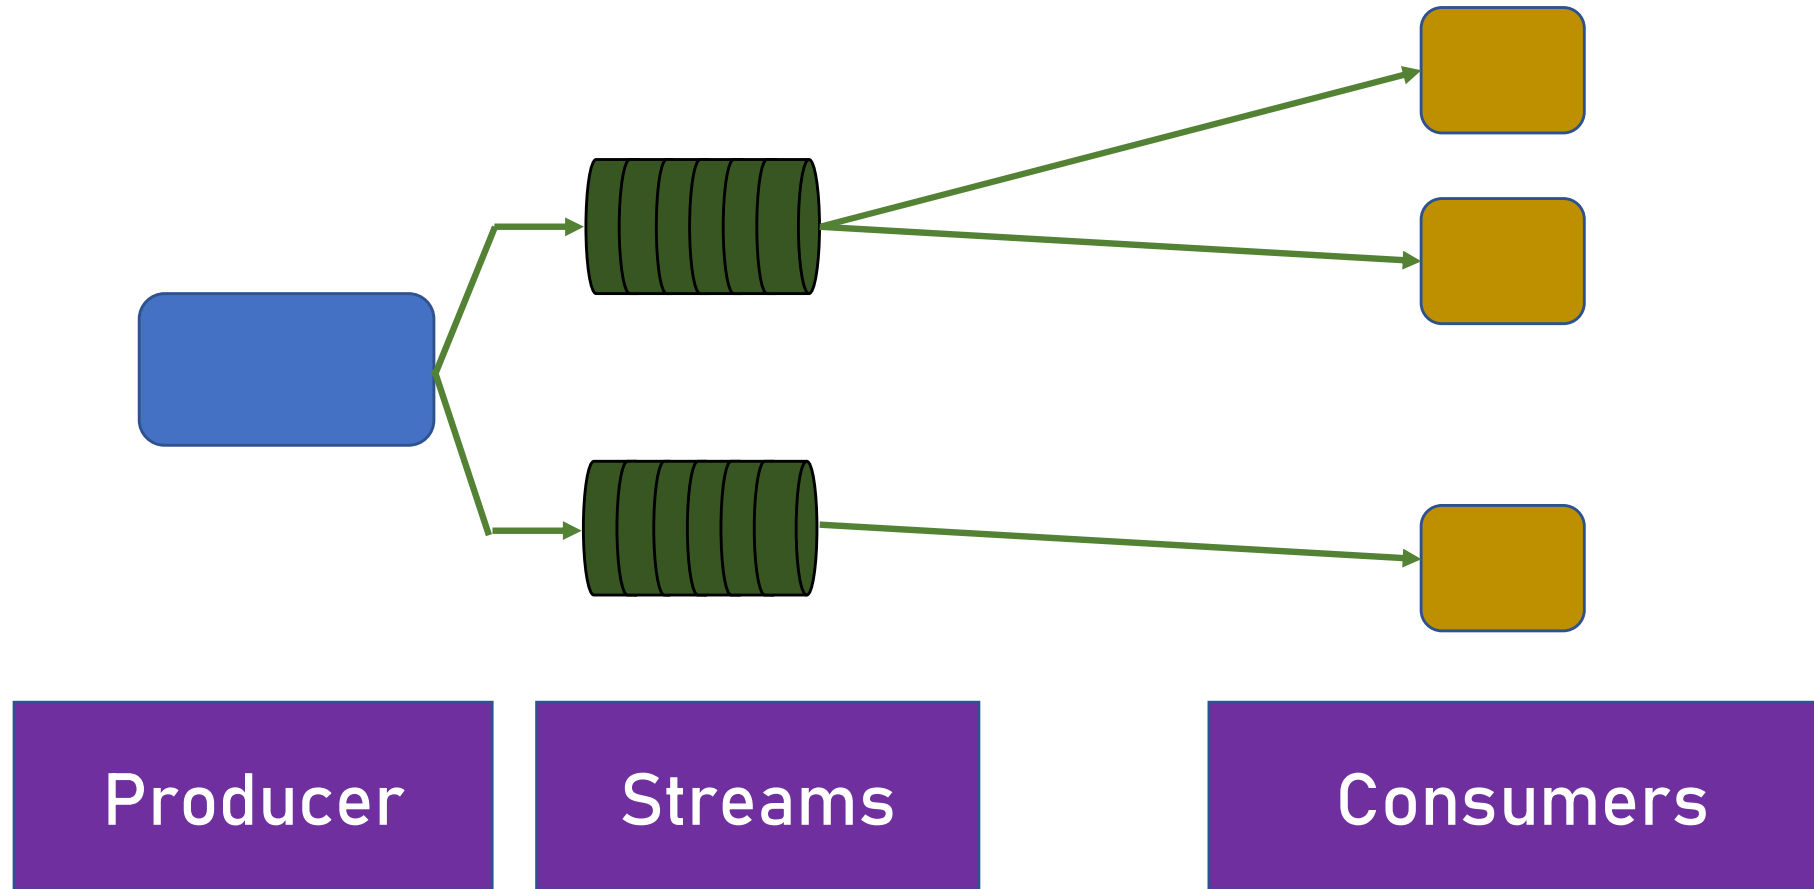
- Event Streaming engines publish stream of events
  - E.g. Telemetry from sensors, system logs etc.
  - The events are published to a “stream”
  - Consumers subscribe to specific stream
  - Events are retained in a stream for a specified amount of time

# Regular EDA





# Event Streaming



# Event Streaming

---

- Consumers can retrieve events that were sent in the past (usually up to a few days)
- Streaming Engine can be used as a central database
  - A single source of truth
- Not all events are necessarily handled
  - Some might be not relevant

# Event Streaming vs EDA

---

## Event Streaming

- Usually used for events generated outside of the system
- Events are retained
- Not all events are handled
- High load

## EDA

- Usually used for events happening inside the system
- Events are not retained
- All events are handled
- No high load

# When to Use Event Streaming

---

- When the system needs to handle stream of events from the outside
  - E.g. Sensor data, logs, etc.
- When events should be retained for future use
- When high load is expected

# Implementing Event Driven Architecture

---

- Mainly 4 things to consider:

Events Approach

Implementing the Channel

Implementing the Producer

Implementing the Consumer

# Events Approach

---

- Two main approaches for implementing events:

Events are retained

Events are not retained

# Retaining Events

---

- The channel retains the event for future handling
- A retention period is defined which after it expires – the event is removed
- Great for streaming events and when the channel is the source of truth

# Not Retaining Events

---

- The channel publishes the events and does not store them
- If a consumer missed an event – it can't be replayed
- Used mainly for in-system events



# Implementing the Channel

- Depends on the events approach

Events are retained

- Use a messaging / queue engine
- Common engines:
  - RabbitMQ
  - Kafka

Events are not retained

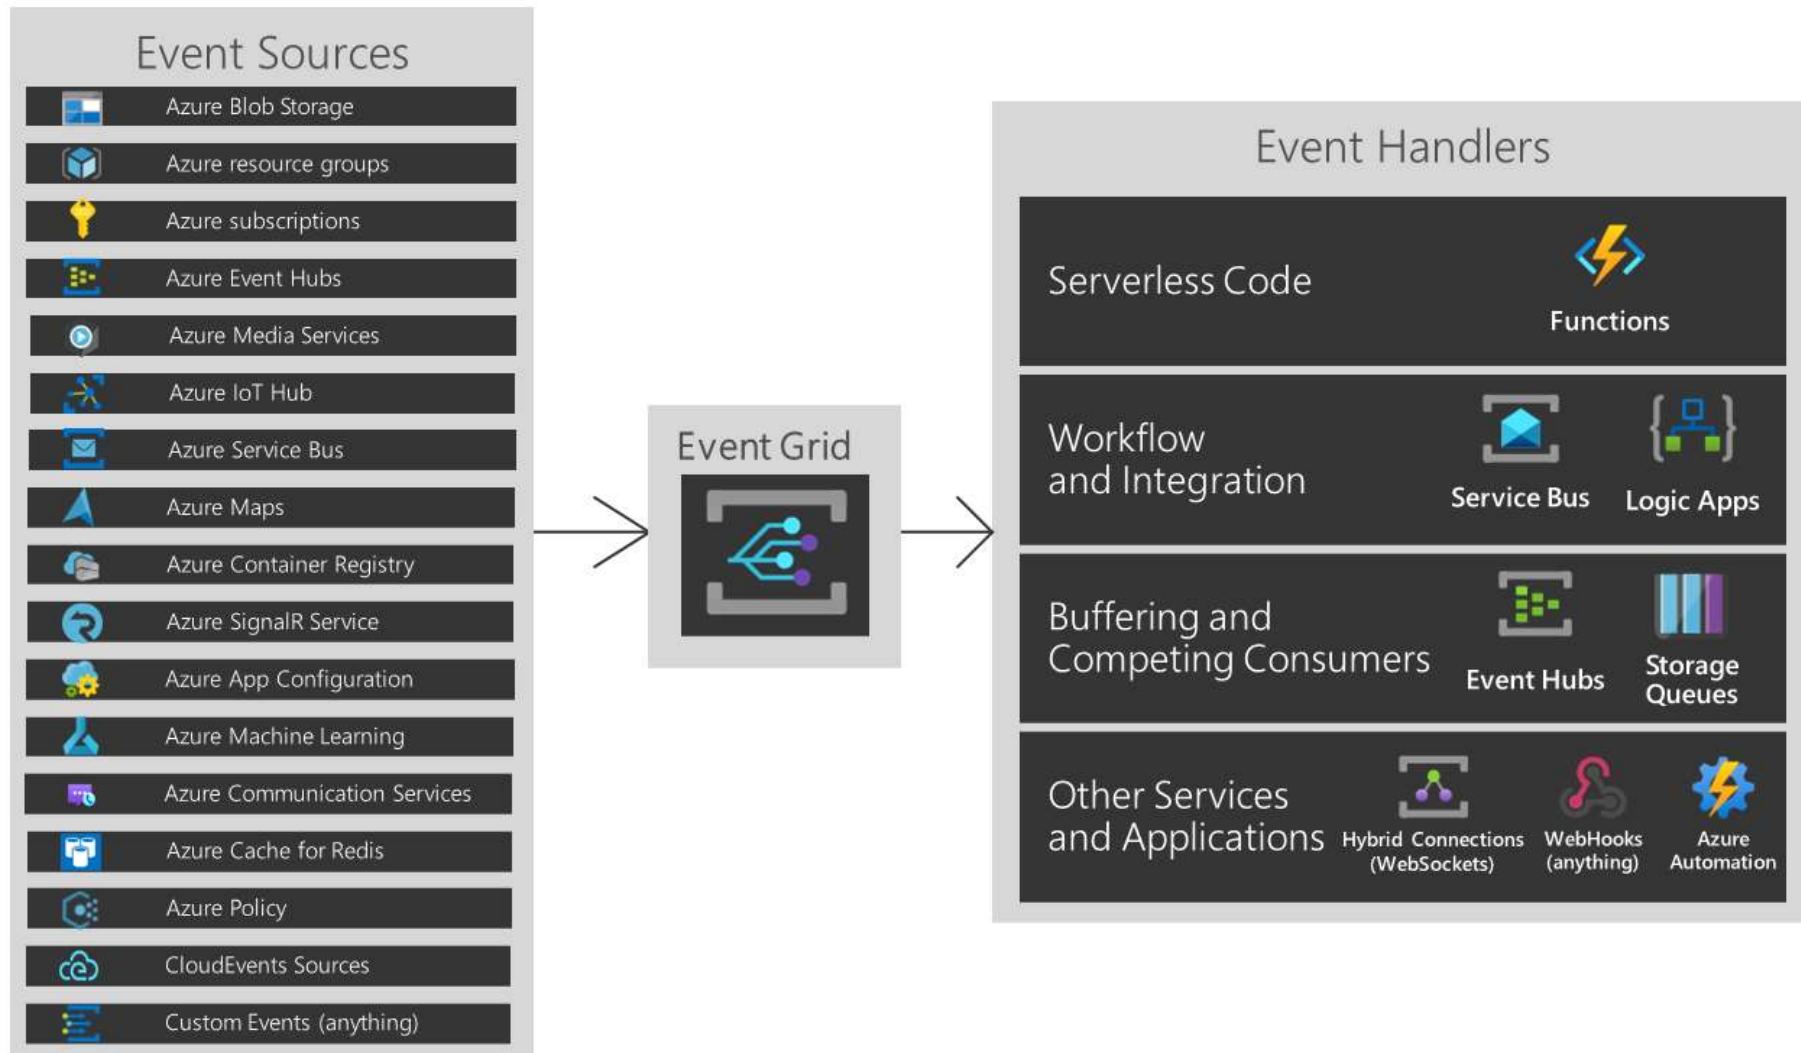
- Use an event publisher
- Specific engine depends on platform used, types of interfaces and more
- Let's see some examples...

# Azure Event Grid

---

- Events publisher in the cloud
- Fully hosted in the Azure cloud, no installation required
- Great integration with a lot of event sources and handlers
- Can deal with thousands of events / sec

# Azure Event Grid



Source: <https://docs.microsoft.com/en-us/azure/event-grid/overview>

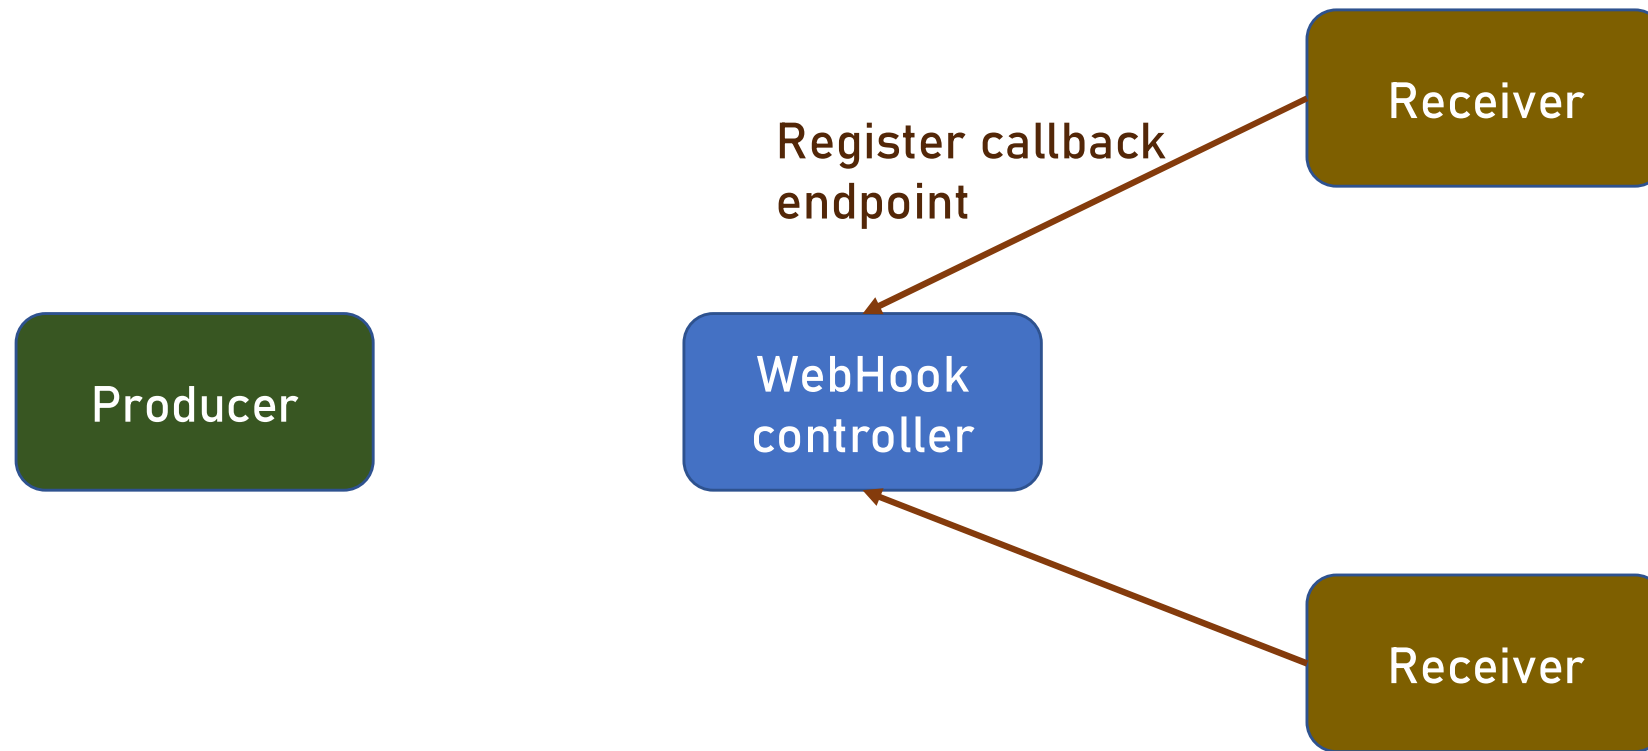
# WebHooks

---

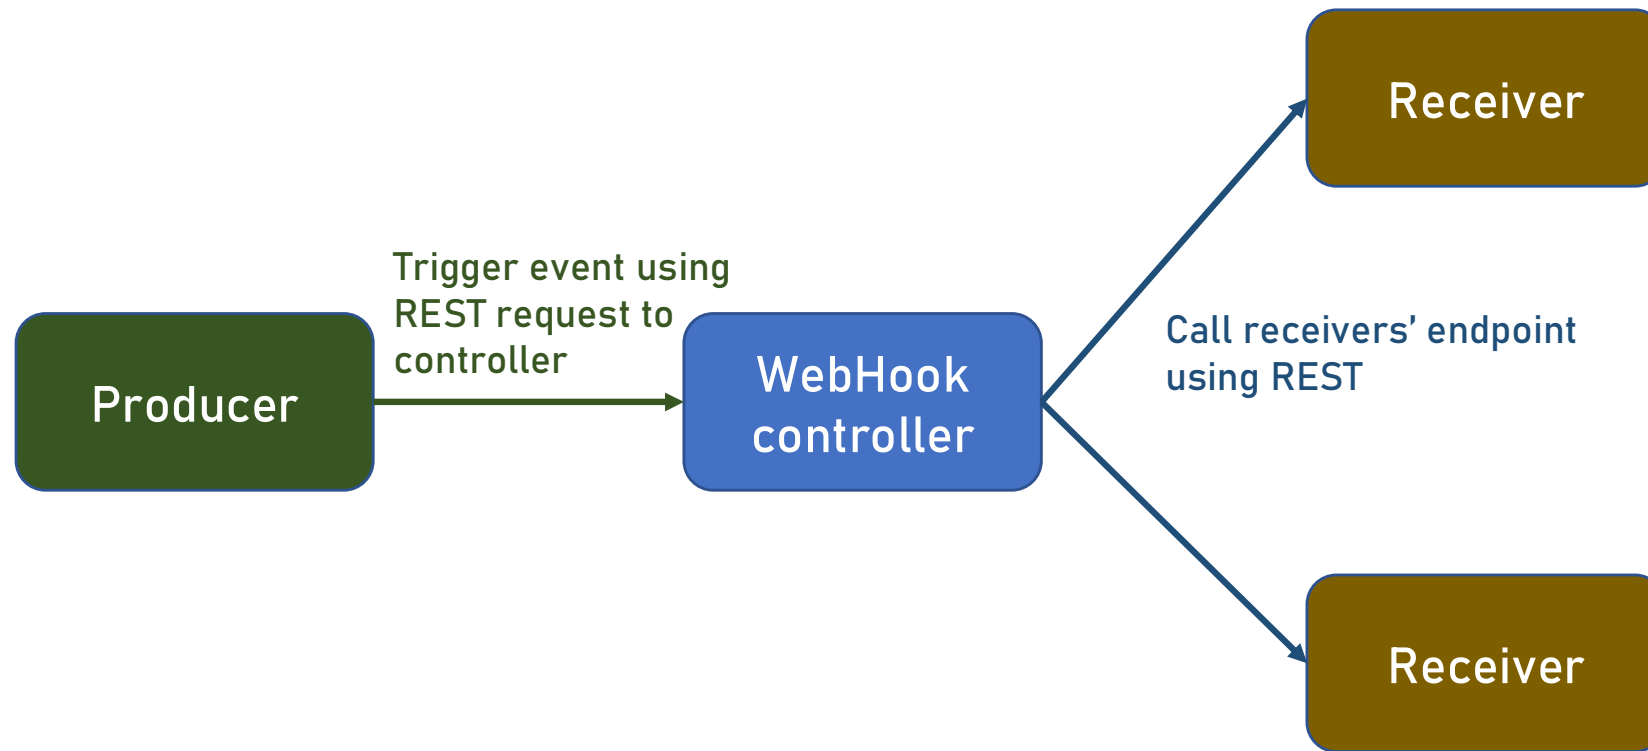
- A standard for publishing events using REST API
- Consumers subscribe to the WebHook engine and register a REST API endpoint that will be called when an event occurs
- The WebHook will call the endpoints when event is triggered
- Easy to implement
- Supported by GitHub, DropBox, PayPal, Stripe and more

# WebHooks Flow

---



# WebHooks Flow



# WebHooks Implementation

---

- Libraries for implementing WebHooks in various platforms
  - E.g. ASP.NET WebHooks
- Websites offering WebHooks:
  - Zapier
  - Ifttt
  - HostedHooks
  - And more...

# HTTP Push Notification


---

- Send events from the server to client(s)
- Great for chats, message notification and more



# HTTP Push Notifications

---



# Implementing Push Notifications

---

- Quite a lot of libraries and frameworks:
  - SignalR
  - Socket.IO
  - gRPC
  - And more...

# How to Choose Event Publisher

| Use...                      | When...  |
|-----------------------------|--|
| Azure Event Grid or similar | <ul style="list-style-type: none"><li>- Hosted in the cloud</li><li>- Need strong integration between backend services</li></ul> |
| WebHooks                    | <ul style="list-style-type: none"><li>- Receivers expose REST API</li><li>- Need something simple and quick</li></ul>            |
| HTTP Push Notification      | <ul style="list-style-type: none"><li>- Need to notify the end user</li></ul>  |

# Remember...

---

**NEVER** develop your  
own channel

# Implementing the Producer

---

- Can be based on any platform
- Needs to be able to communicate with the channel
- Depends on the channel implementation
- Let's see some examples...

# Implementing the Producer

---

- RabbitMQ
  - Use the RabbitMQ client library for your platform
  - There's one for almost every platform
  - Choose from:
    - <https://www.rabbitmq.com/devtools.html>

# Implementing the Producer

---

- SignalR
  - Install the SignalR library
  - Configure the Hub
  - Allow connections from clients
  - Define functions that will send messages to clients
  - Optional – create groups to filter messaging

# Implementing the Consumer

---

- Can be based on any platform
- Needs to be able to communicate with the channel
- Depends on the channel implementation
- Let's see some examples...



# Implementing the Consumer

---

- RabbitMQ
  - Use the RabbitMQ client library for your platform
  - There's one for almost every platform
  - Choose from:
    - <https://www.rabbitmq.com/devtools.html>

# Implementing the Consumer

---

- WebHooks
  - Register the consumer using the WebHook REST API
  - Expose REST API that will be called by the WebHooks

# Our System

---

- Introducing:

**NOP**

The N0ise Processing system

# NOP

---

- A system for receiving and processing noise data from external sensors
- The system should:
  - Receive the telemetry
  - Validate it
  - Notify clients on new data

# NOP

---

- The data is a number representing the decibels recorded
- Every sensor sends the data every 30 secs
- ...That means that if there are a lot of sensors, there's quite a lot of data...
- E.g. 1000 sensors => 33 msgs / sec

# NOP Design Requirements



Handle load

Streaming engine should be used.  
Processors pull from the stream  
when possible

Validate the data

The first thing that should happen  
after receiving the data

Unknown number of clients

Classic events requirement

No sync users' commands

No synchronous actions required

# NOP Event Driven Architecture

